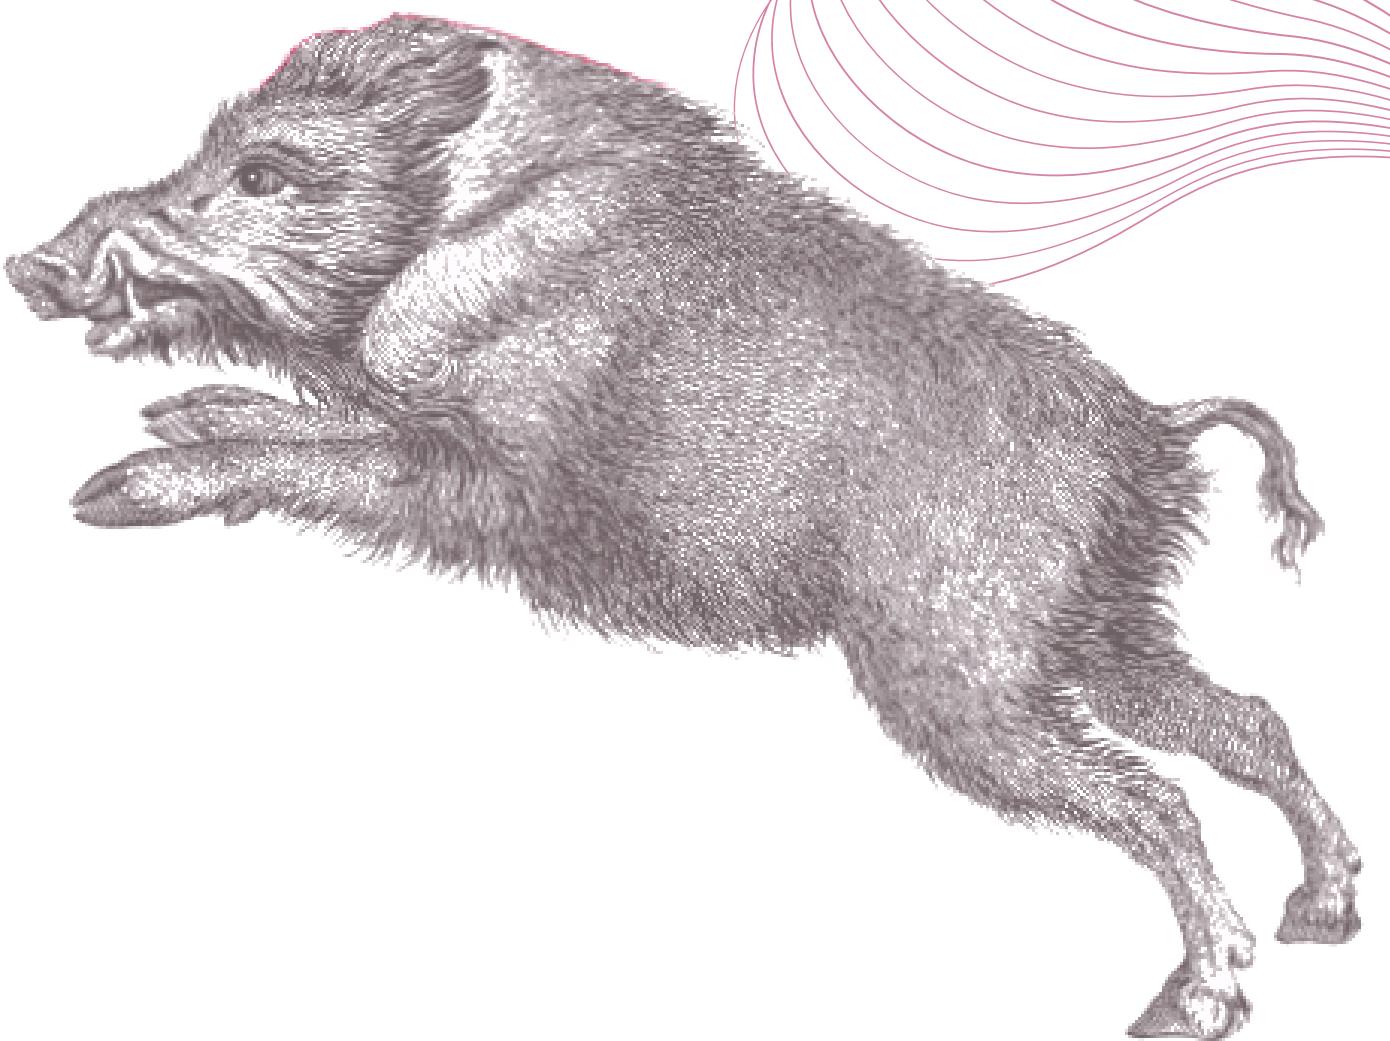
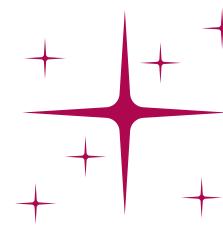


Designing Data-Intensive Applications

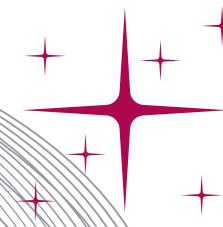
A book by:
Martin Kleppmann



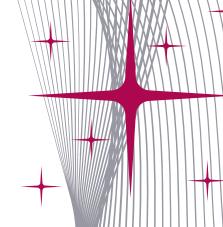
Martin Kleppman



Professor at the University of Cambridge



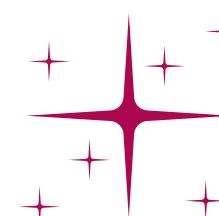
Published Designing Data-Intensive
Applications 2017.



Martin Kleppman co-founded Rapportive
(acquired by LinkedIn in 2012) and Go Test
It (acquired by Red Gate Software in 2009)



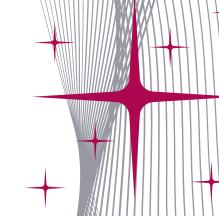
Book Introduction



Data intensive >>>>> Compute Intensive



Data systems are such a successful abstraction that we use them all of the time without thinking about it



This book will talk about the principles and practicalities of data systems.

O'REILLY®

Designing Data-Intensive Applications

THE BIG IDEAS BEHIND RELIABLE, SCALABLE,
AND MAINTAINABLE SYSTEMS



Martin Kleppmann

Data Intensive App Example



Why It's Data-Intensive?

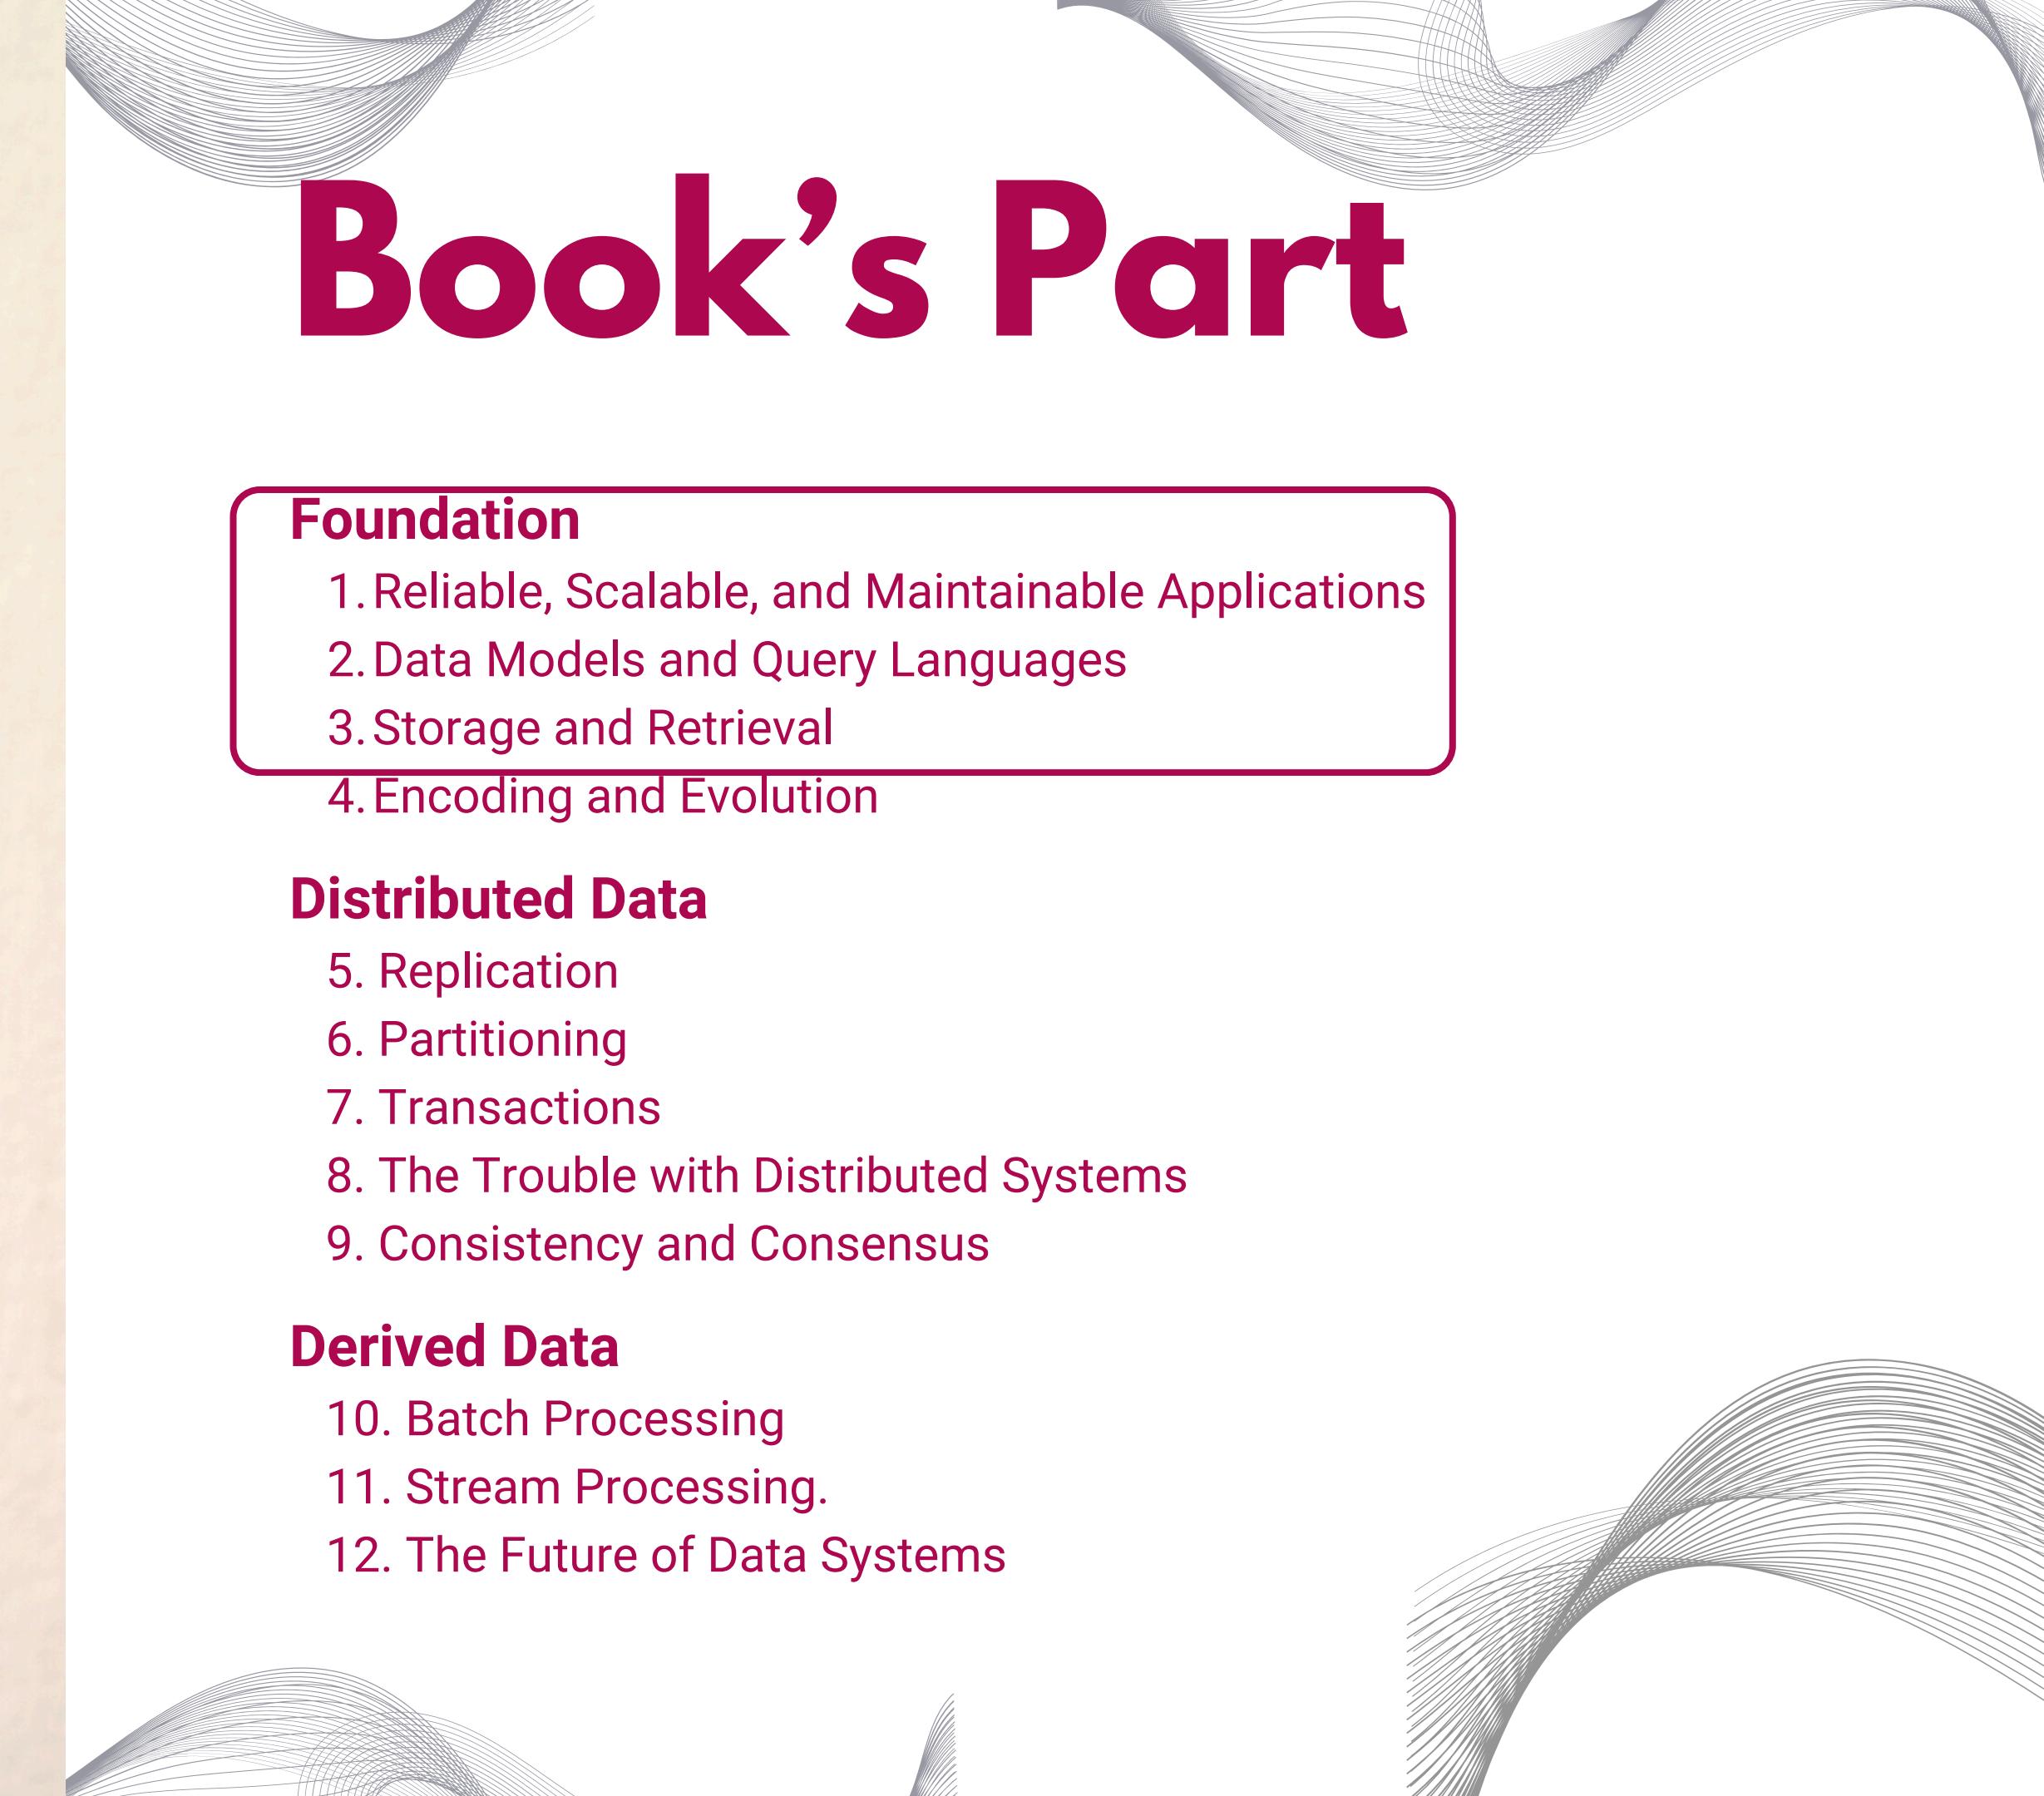
- Massive Data Volume:** Processes millions of taxpayer records, transactions, and documents daily.
- Complex Processes:** Validates ever-changing tax rules with intricate data relationships.
- Distributed System Complexity:** Relies on many interconnected services (e.g., banks, regional tax offices, external services) that must stay continuously synchronized.

What happened?

- Service Disruption:** Frequent system crashes
- Hybrid System Rollback:** Due to instability, they reverted to a parallel system with the old tax platform.
- Financial & Public Trust Loss:** Despite a Rp 1.3 trillion budget, issues persist, leading to extra costs and eroding public confidence.



Book's Part



Foundation

- 1. Reliable, Scalable, and Maintainable Applications
- 2. Data Models and Query Languages
- 3. Storage and Retrieval
- 4. Encoding and Evolution

Distributed Data

- 5. Replication
- 6. Partitioning
- 7. Transactions
- 8. The Trouble with Distributed Systems
- 9. Consistency and Consensus

Derived Data

- 10. Batch Processing
- 11. Stream Processing.
- 12. The Future of Data Systems

Chapter 1



Reliable, Scalable, and Maintainable Applications

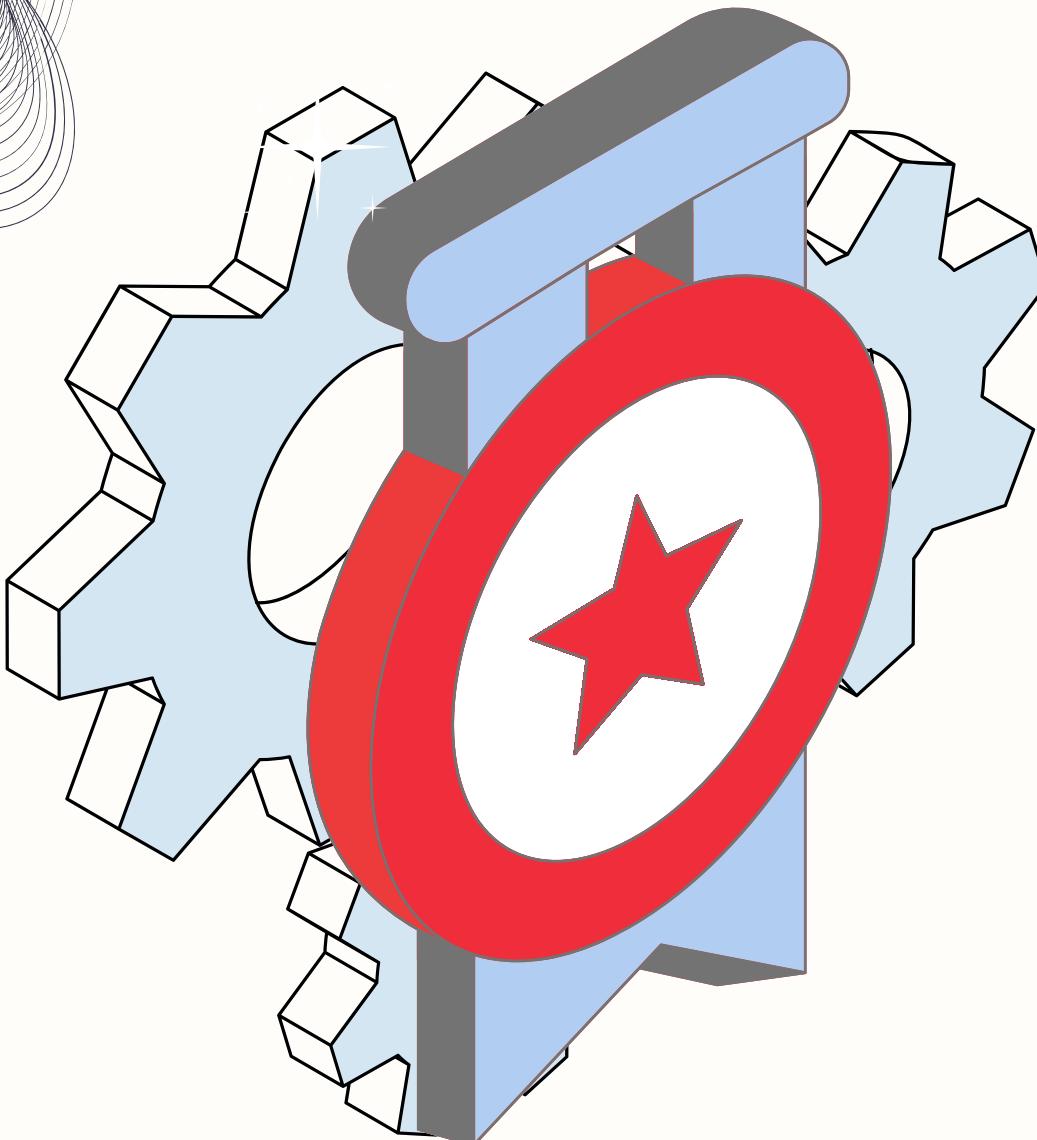
Reliability

Definition

Continuing to work correctly even when things go wrong

Why???

To find a solution or an alternative when something goes wrong



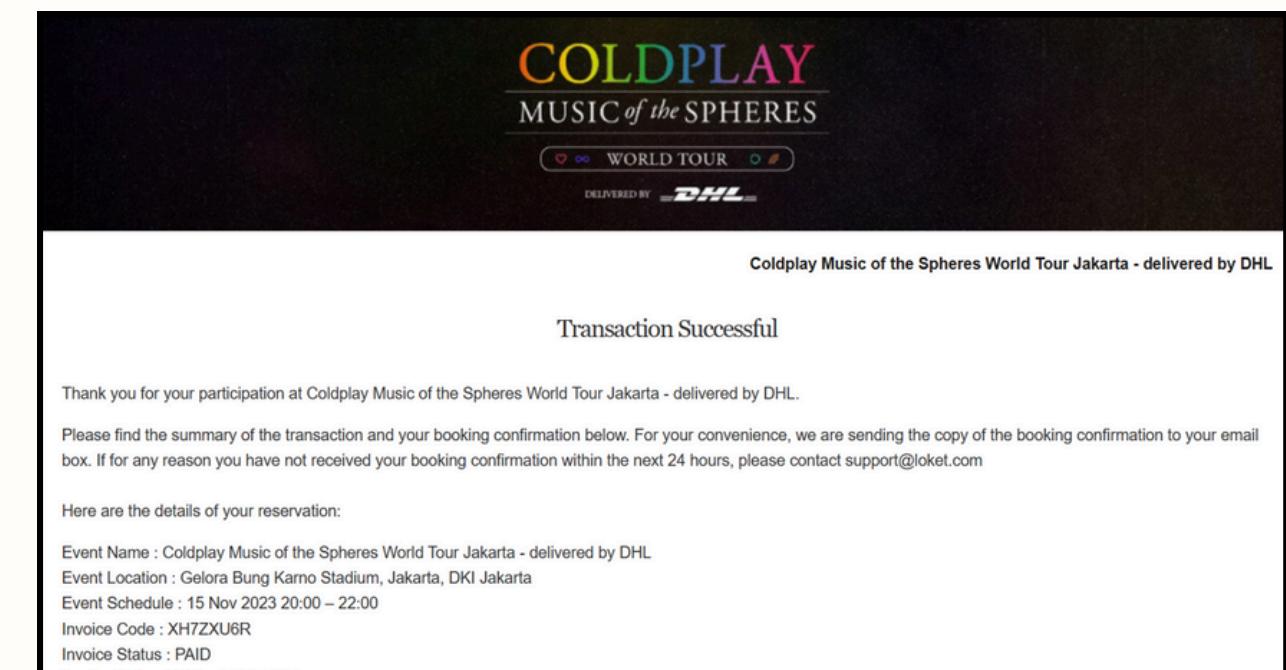
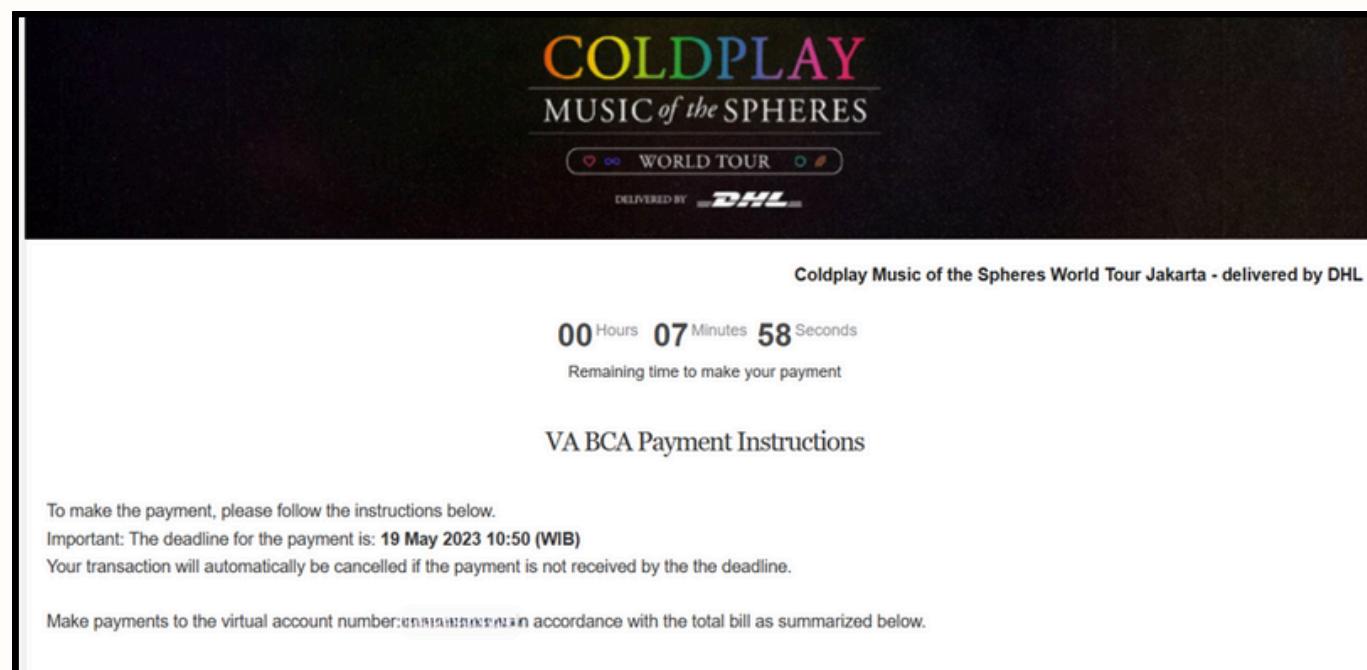
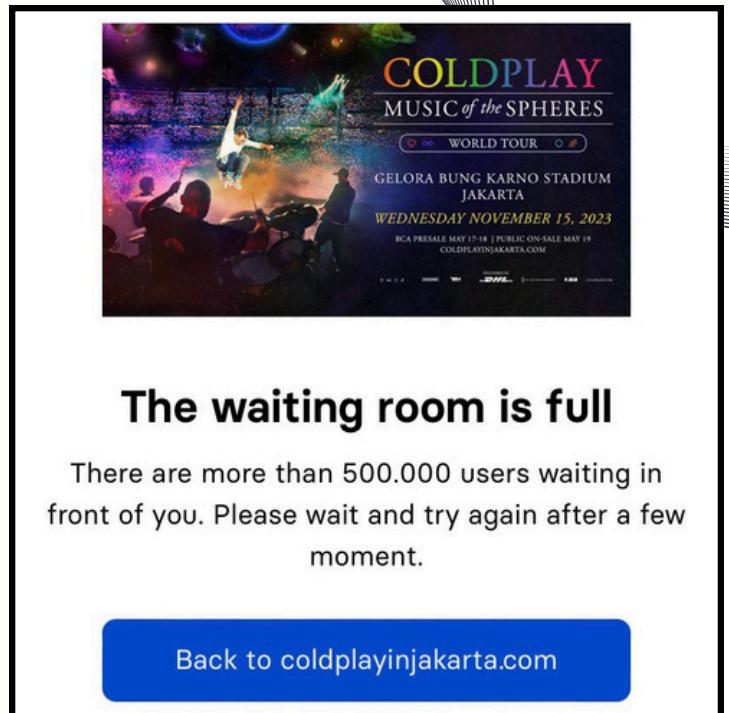
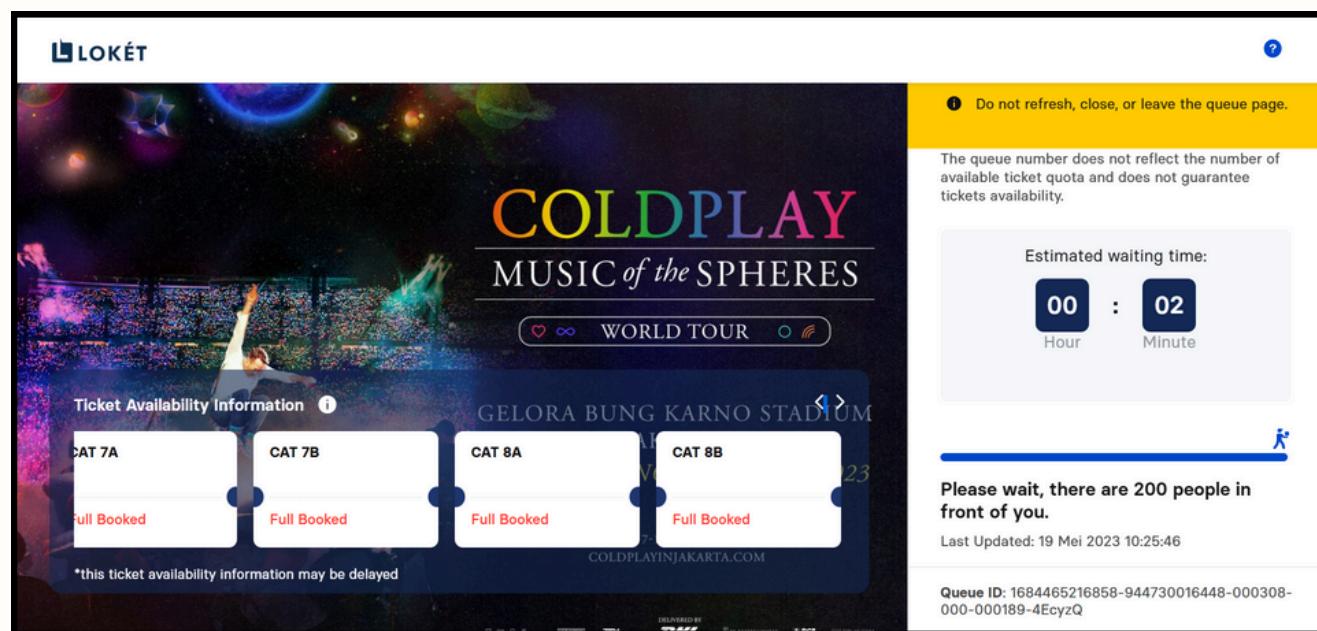
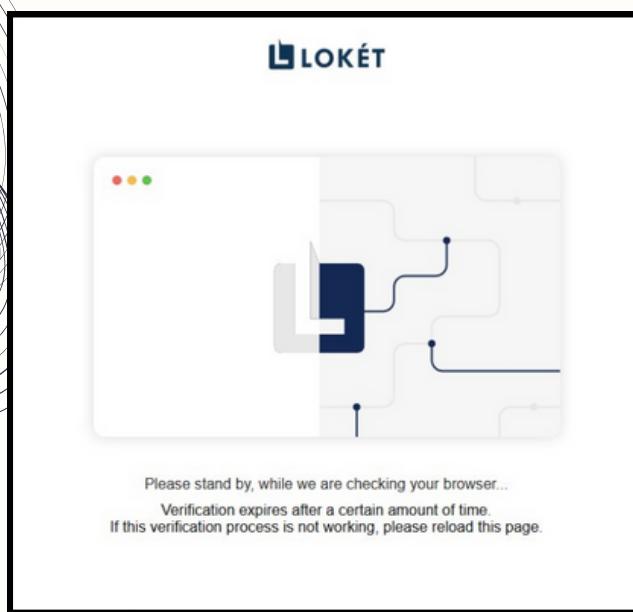
Reliability



Typical Expectation

- Application performs the function the user expected
- Tolerate the user making mistakes or using the software in unexpected way
- Its performance is good
- The system prevents unauthorized access and abuse.

Examples of Good Reliability



Kinds of Faults



Hardware Faults



Software Faults



Human Errors

Hardware Faults

Physical failures in hardware that can disrupt the system.
We can improve resilience to hardware failures by using hardware redundancy (RAID, dual power supplies, hot-swappable components)

Example with hardware redundancy

Initial State: Server has 2 PSUs, one connected to the main power, the other to a UPS/generator.

Issue Occurs: PSU 1 fails, but the server keeps running on PSU 2

Action Taken: Team gets a notification and replaces PSU 1 without shutting down the server (hot-swap).

Recovery: New PSU is installed, and the server returns to using both power sources.

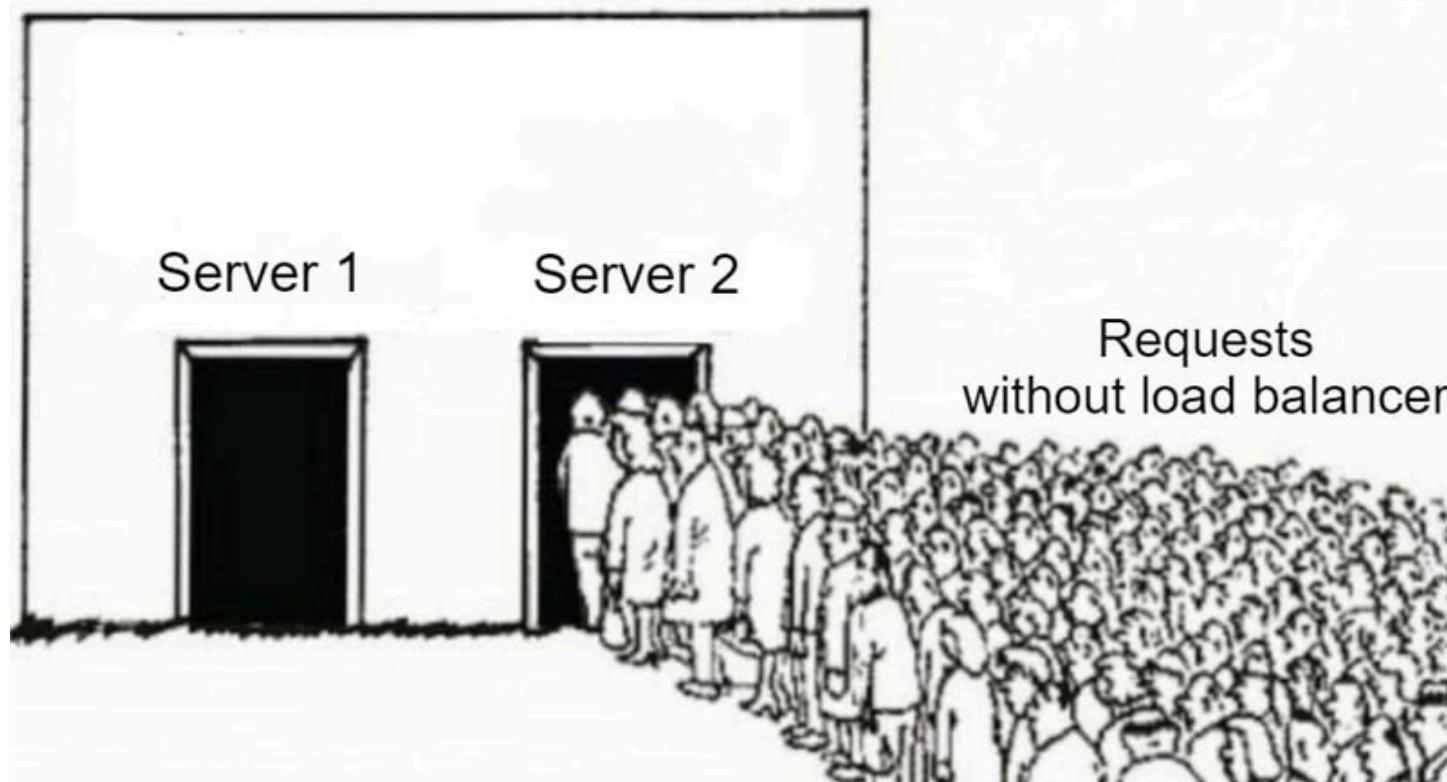
Result: No downtime, data remains safe, and services



Hardware Faults

Software fault tolerance techniques

- Modern systems adopt software fault-tolerance techniques (data replication, load balancing, rolling upgrades)



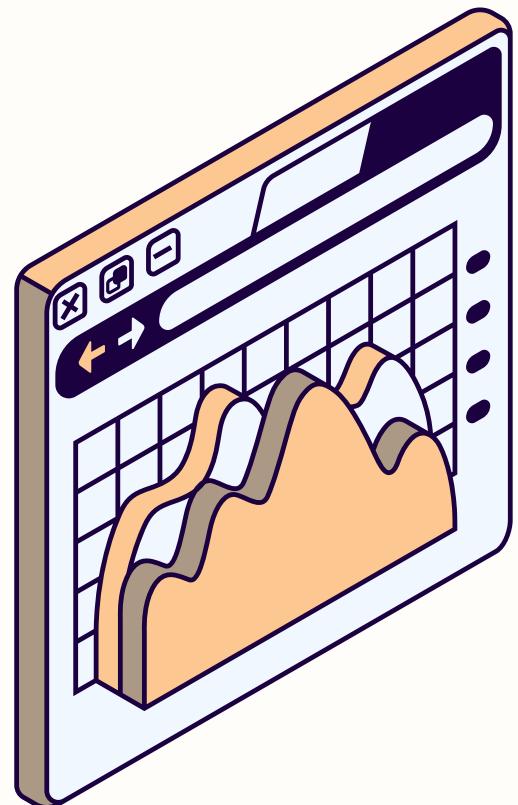
Example

- **Initial State:** The app runs on 3 servers with a load balancer.
- **Issue:** 1 server crashes, but the other 2 handle requests.
- **Action:** The load balancer redirects traffic to the active servers.
- **Recovery:** After repair, the server is reintegrated automatically.
- **Result:** Zero downtime, uninterrupted service, and system stability.

Software Faults

Characteristic

- Difficult to detect
- Systematic & cascading failures
- No quick fix



Example

- **Initial State:** A video streaming app with multiple microservices.
- **Issue:** A memory leak in the recommendation service spikes CPU and RAM, disrupting playback.
- **Action:** Engineers isolate, fix, test, and redeploy the service.
- **Recovery:** The service is gradually restored with performance monitoring.
- **Result:** Core services stay online, downtime is minimized, and the bug is resolved.



Tolerating Faults

Human Errors

Causes

- Misconfiguration settings
- Lack of Understanding
- Pressure & Fatigue

Solution

- Minimize Opportunities for Mistakes
- Decouple High-Risk Actions
- Enable Quick & Easy Recovery
- Use Detailed & Clear Monitoring
- Provide Good Management & Training

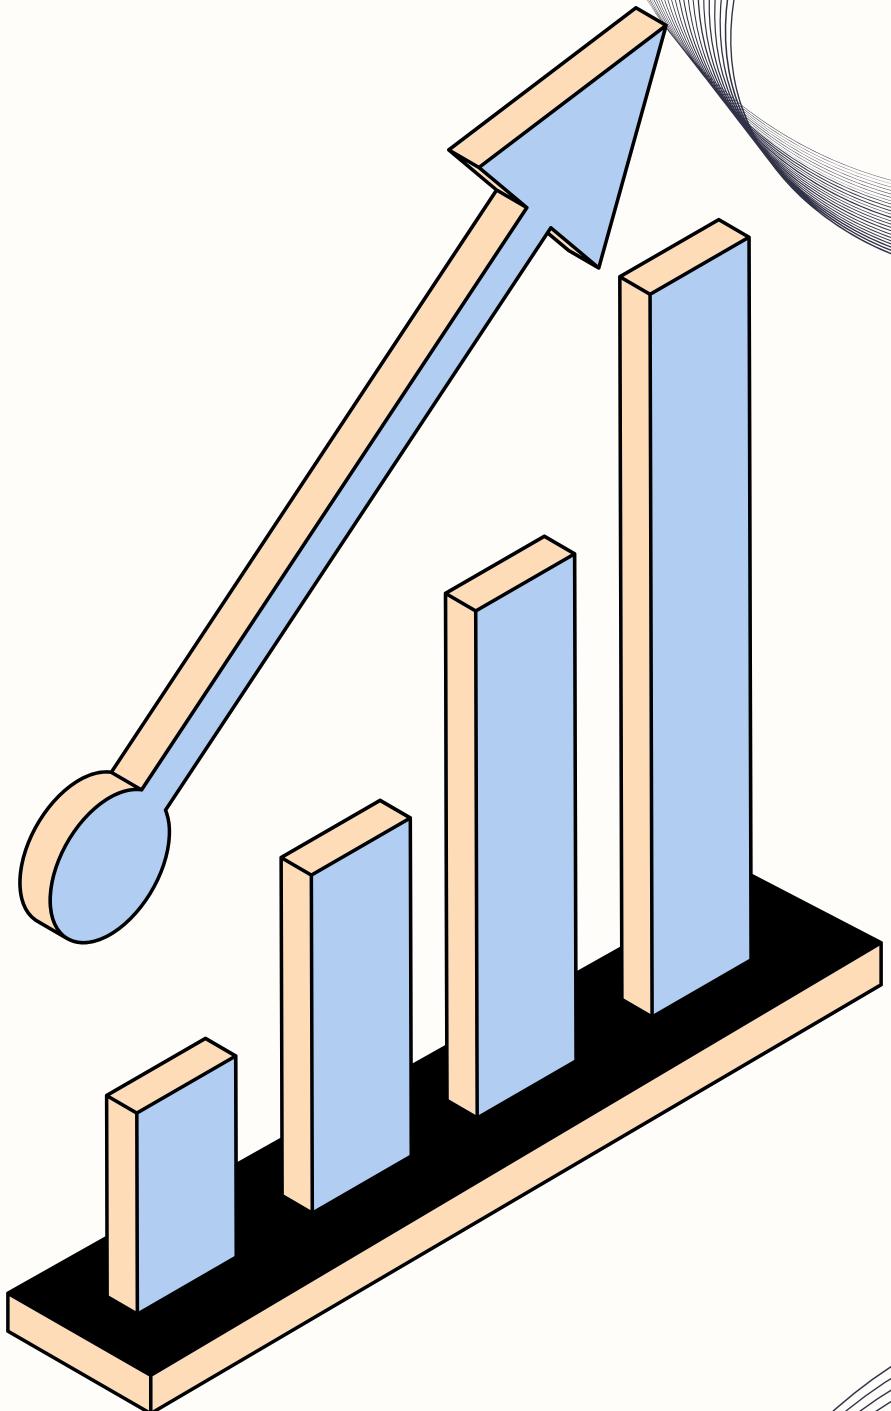
Scalability

Definition

Asking if the system grows in a particular way,
what are the options for coping with that growth

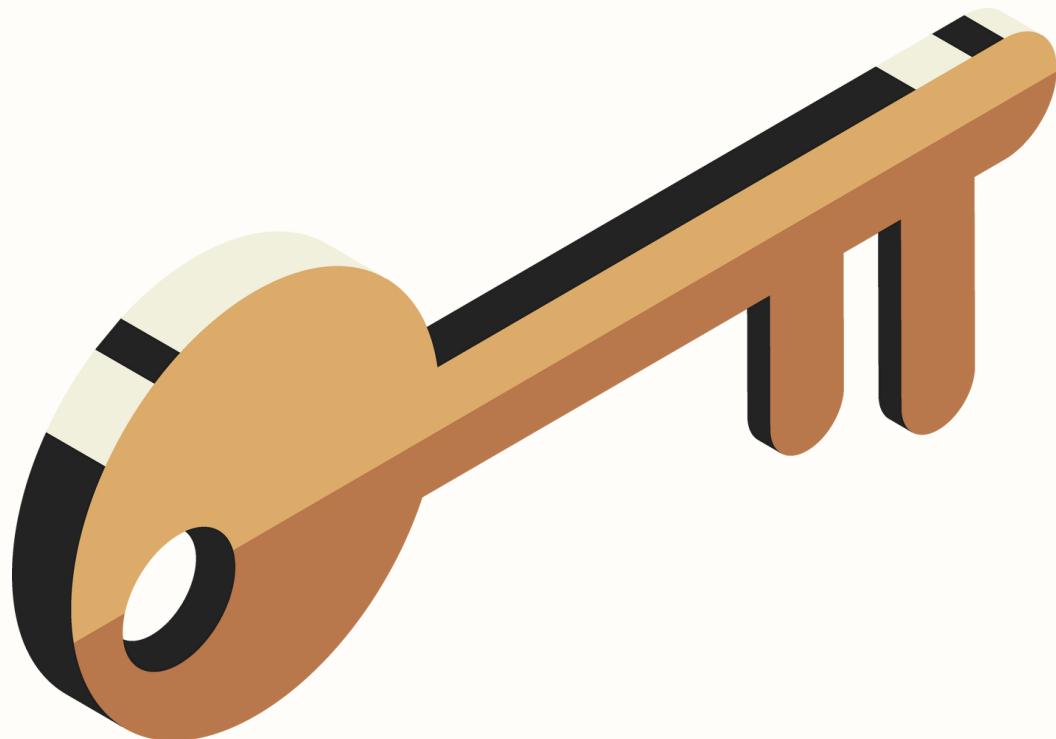
Why???

A system should be able to handle additional loads, such as an increasing number of users or data volume, without experiencing a decrease in performance



Scalability

Key Considerations



- Load
- Performance
- How to cope with load?

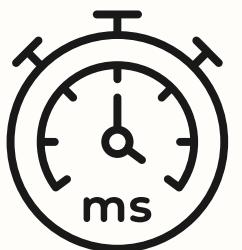
Scalability

Load

- Requests per second
- Read/Write ratio
- Active users
- Cache hit rate

Scalability

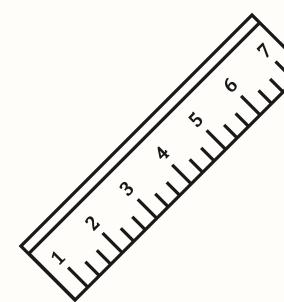
Perfomance



Latency



Response time



Perfomance measured
by percentiles



Percentiles are used in SLA

Scalability

How to Cope with Load?

Scale Up

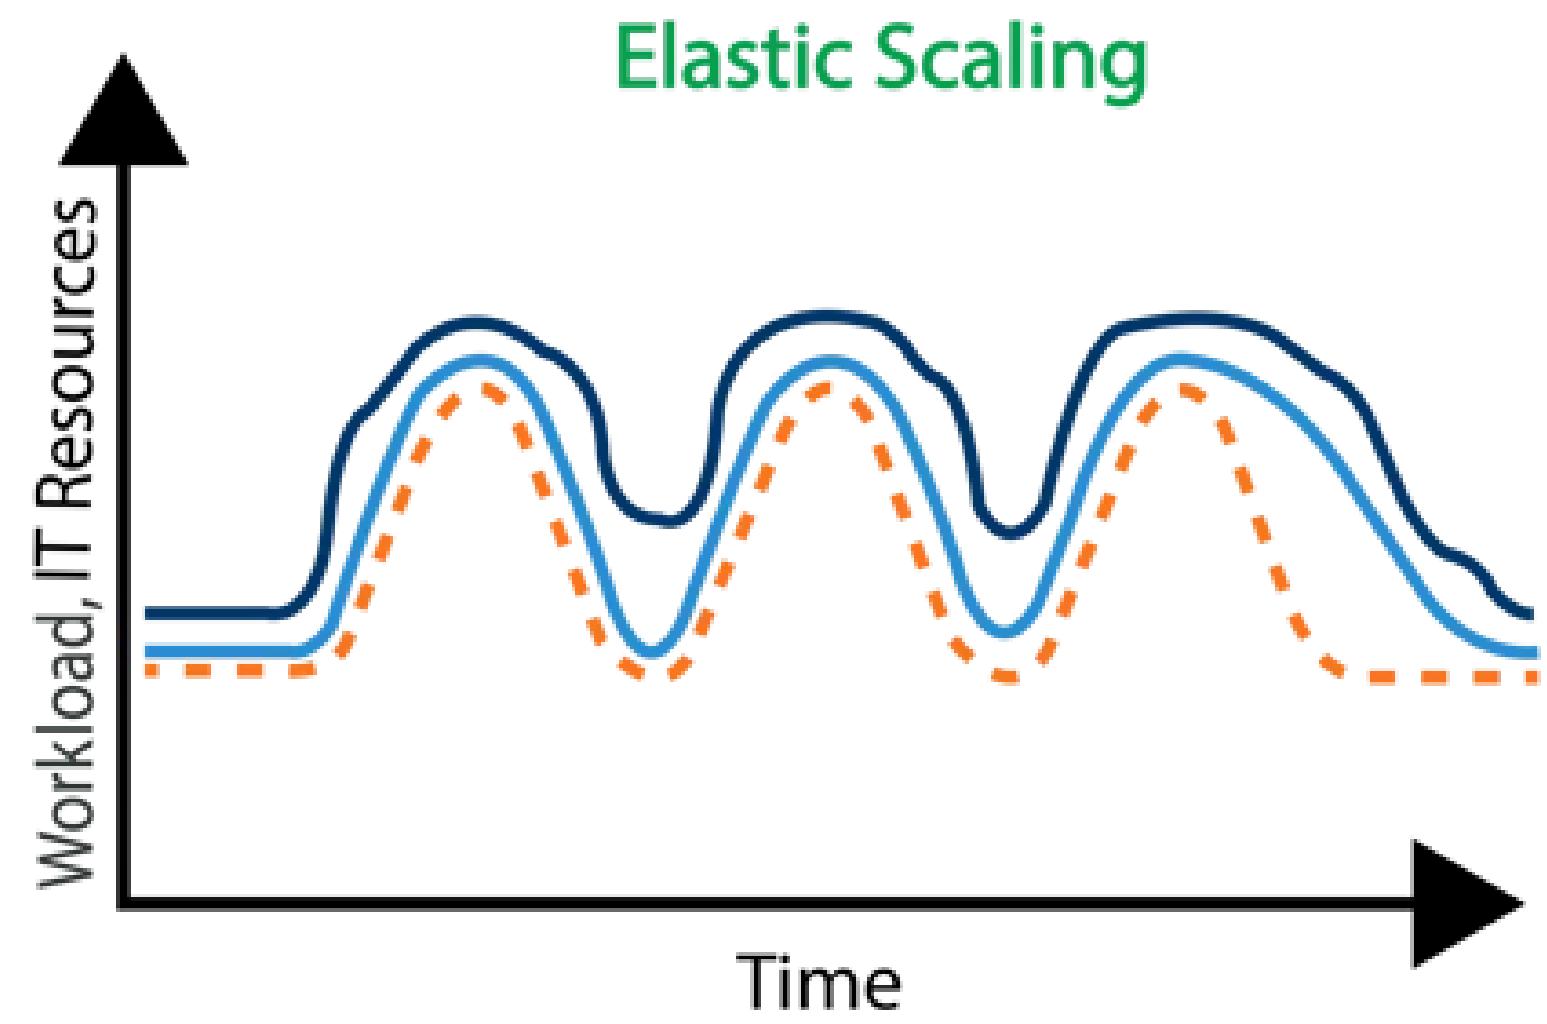


Scale Out

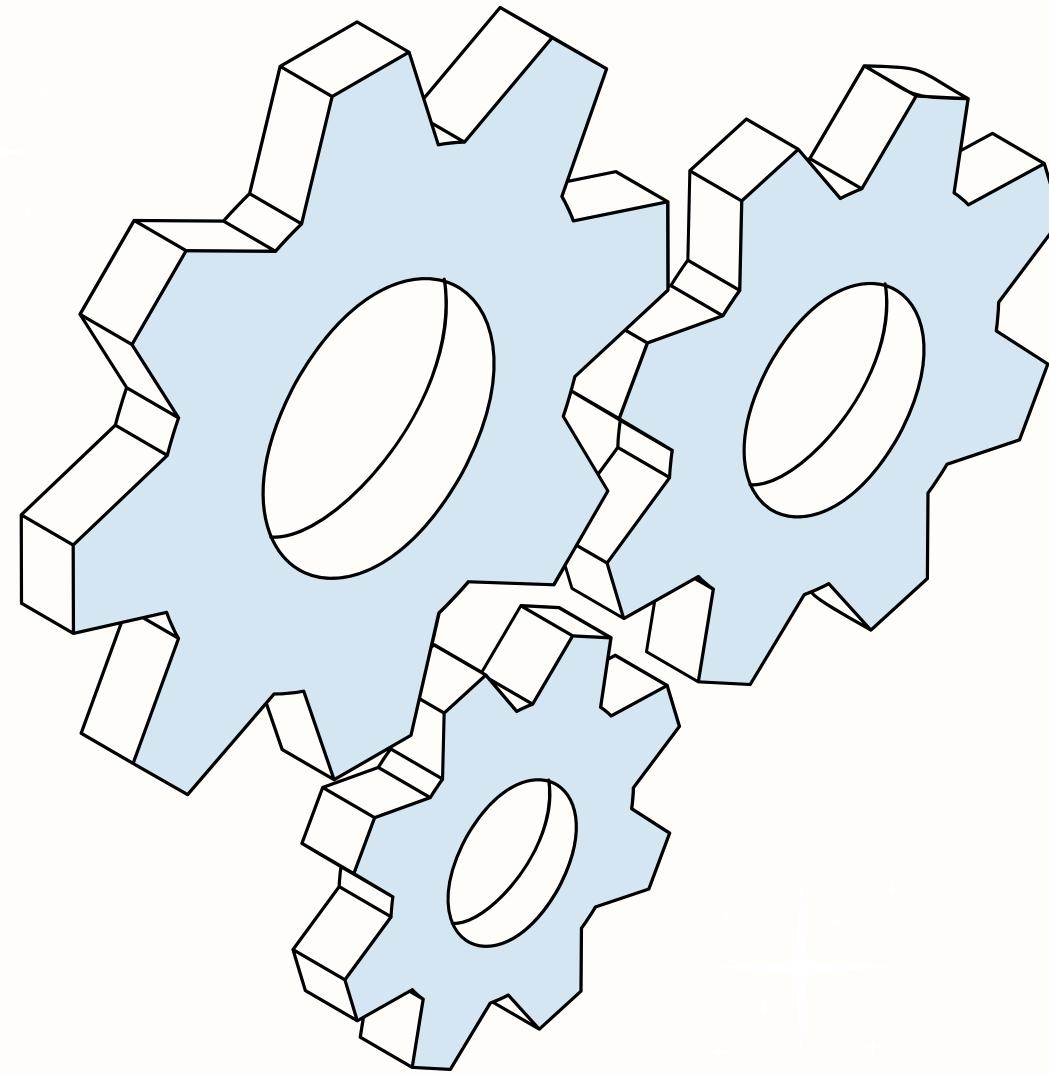


Scalability

How to Cope with Load?



Maintainability



Definition

Keeping the system easy to modify and adaptable for future changes or improvements

Why???

Eases the burden on engineers who work with the system by preventing difficulties in debugging, developing new features, and fixing bugs, which can lead to higher operational costs

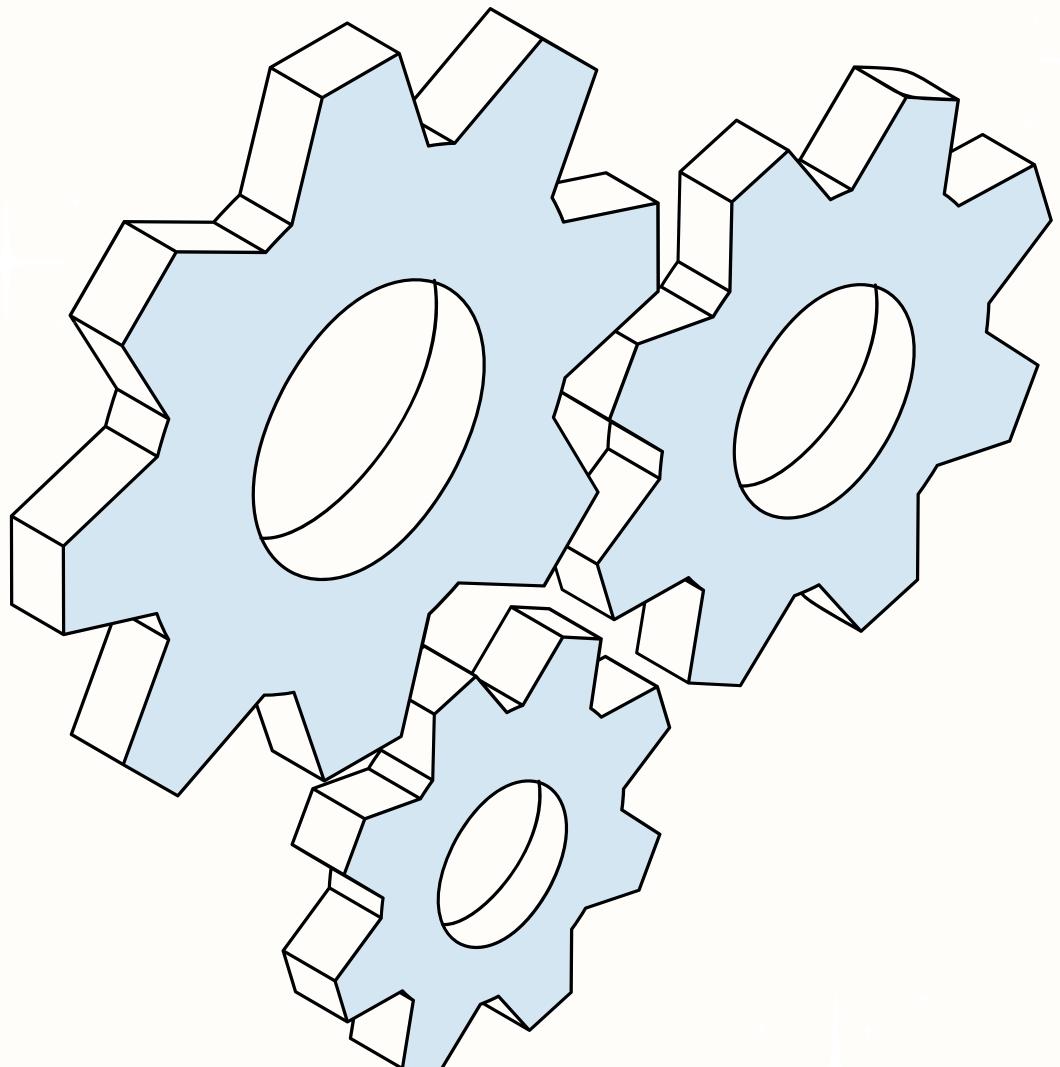
Maintainability



Lack of Maintainability Example

- Initial State: An e-commerce app lacks clear documentation and organized code.
- Issue: Adding a seasonal discount feature is difficult due to overlapping functions and inconsistent variables.
- Impact: Slower development, more bugs, and longer debugging times.
- Result: Delayed feature release, lost revenue, and higher costs to fix the code.

Maintainability



Operability

Simplicity

Evolvability

Operability

Measuring how easy it is for the IT team to operate and maintain the system on a daily basis



How?

- Monitoring system and quickly restoring
 - Tracking system failure and degraded performance
 - Keeping software up to date
 - Keeping track how one system affect the other
 - Establishing good configurations practices
 - Avoiding dependency or individual machines
 - Good documentation
- deployment,
individual

Simplicity

What is simplicity?

- The simpler the design and code, the easier it is engineers to understand



How to reduce complexity

- Clean Code
- Modular Design
- Use descriptive and consistent names
- Refactoring
- Using design patterns
- Documentation & comments

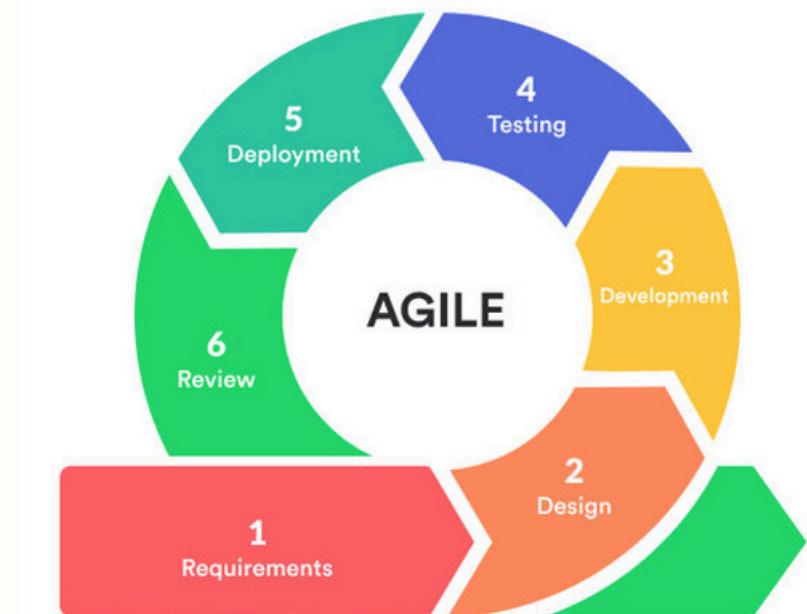
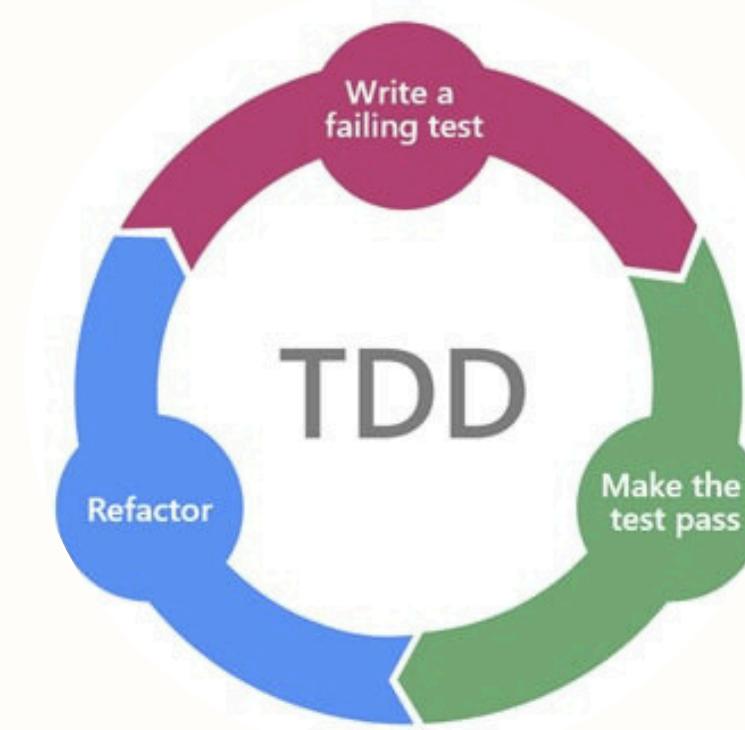
Evolvability

Making it easy for engineers to make changes into the system in the failure

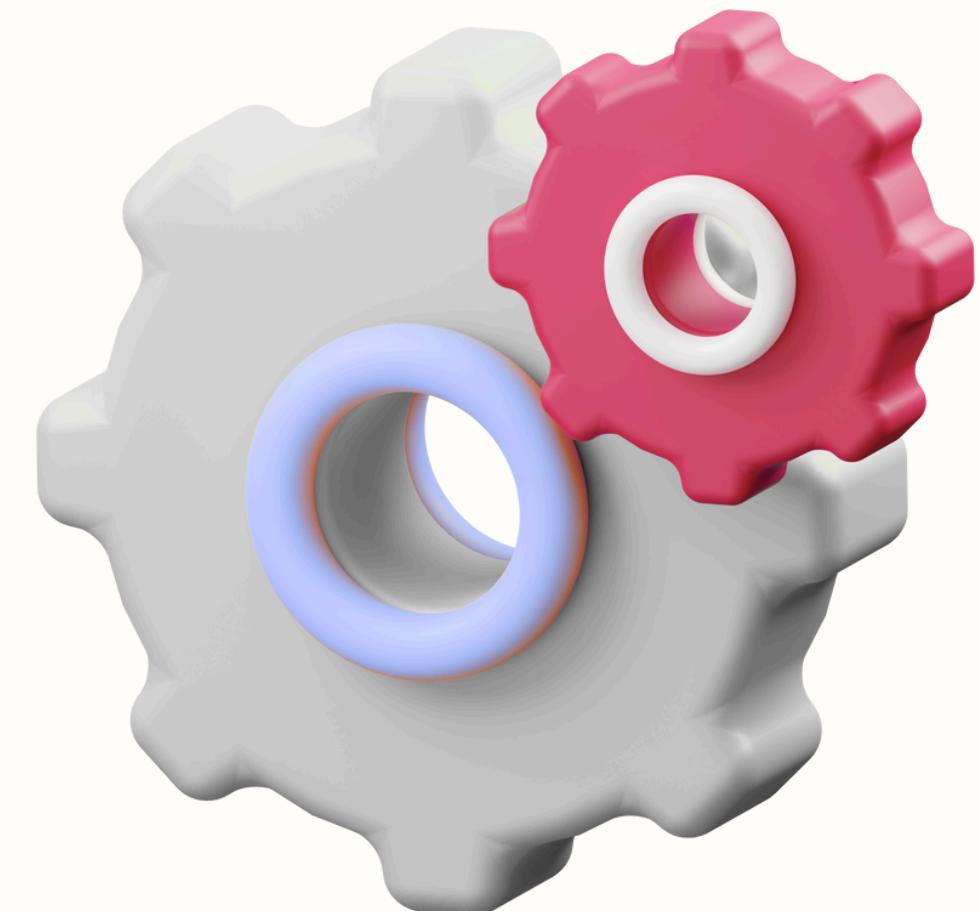
Why it is needed?

- Business Priorities change
- Users request new features
- New platforms
- Architectural changes

How?



Chapter 2



Data Models and Query Languages

Data Models

Definition

A way to represent data at a level that is easier for humans and systems to understand

Importance

Data models shape how we think about problems, influence how software is designed, and determine which operations are efficient or challenging

Abstraction

**APIs/Data structures/Objects
at application layer**

**Document or Relational
Database**

**Bytes in
memory or
disk**

H/W

Data Models

Relational

Document

Graph

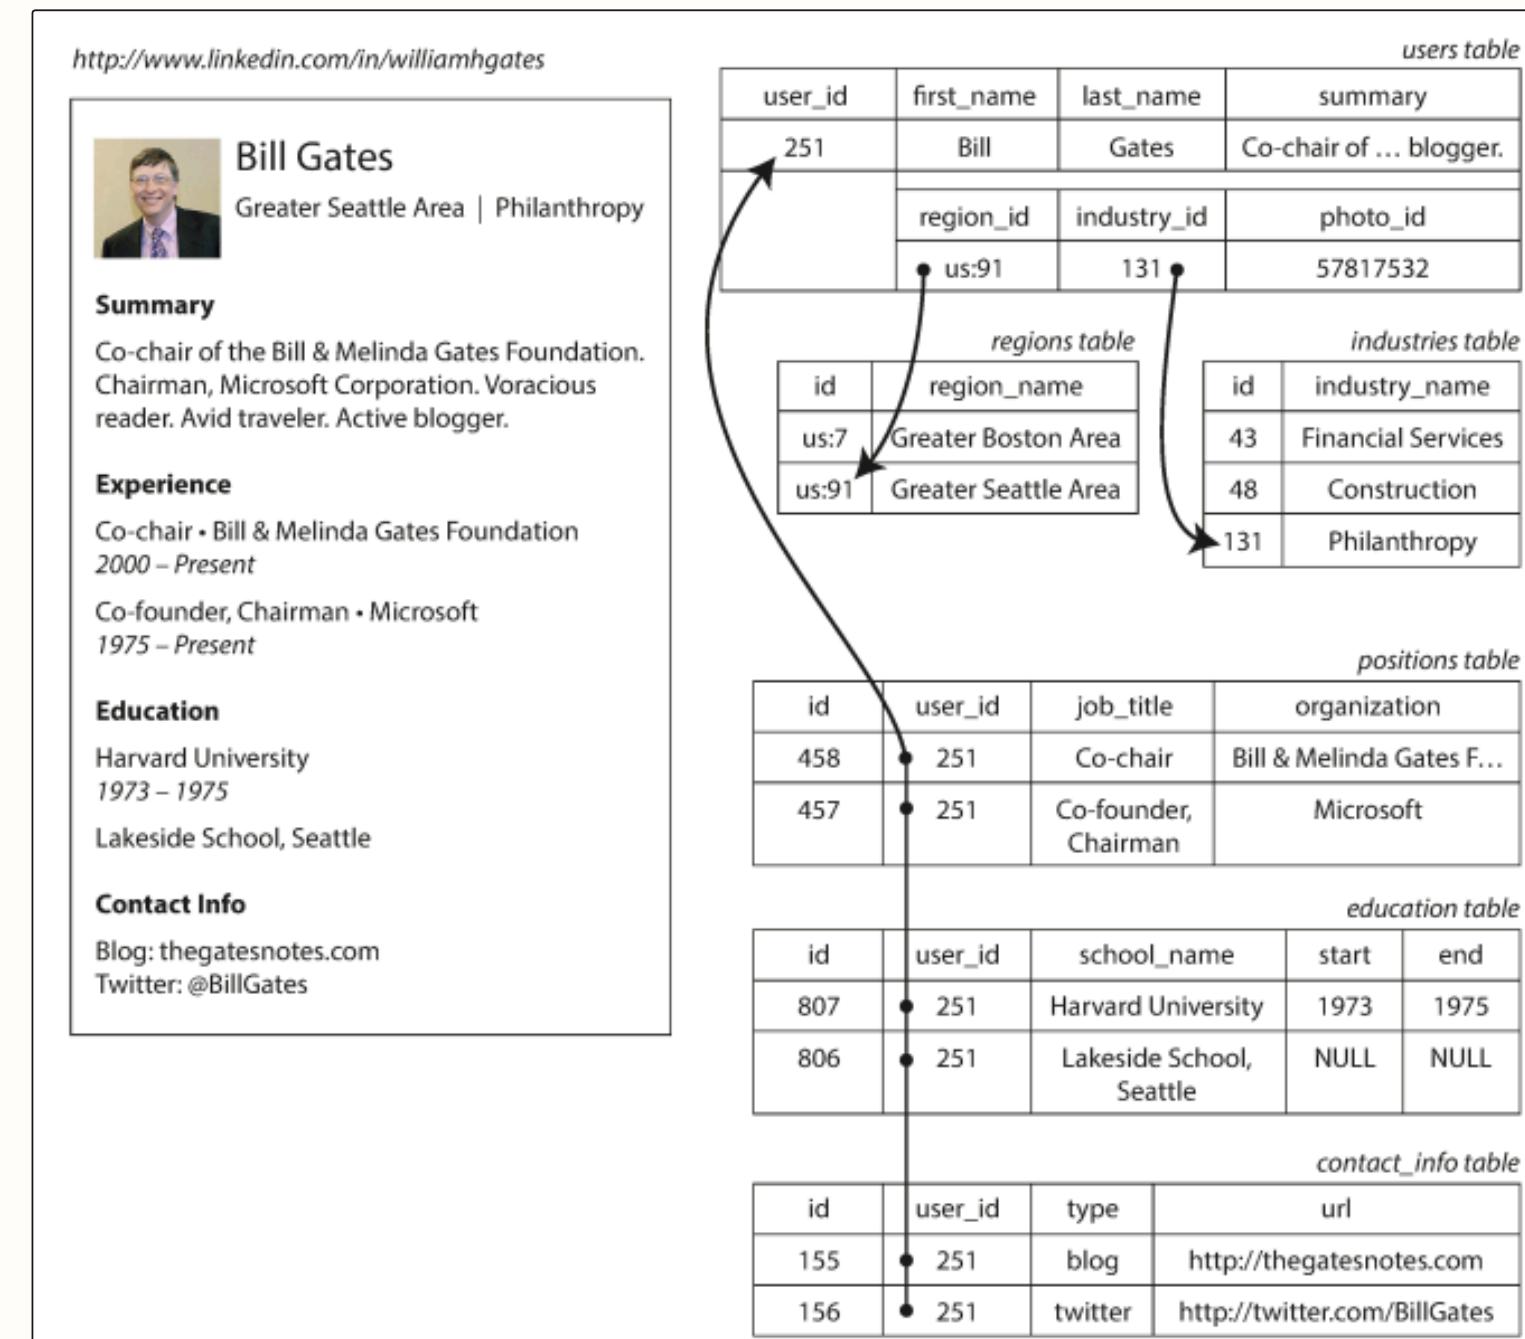
Relational

Definition

- A relational database organizes data into relations, also known as tables, where each relation is an unordered collection of tuples (rows).
- Schema on write

Problem

- Impedance mismatch
- Limited Scalability
- Rigid Structure
- Excessive Joins



Document

NoSQL

- A type of NoSQL database that stores, retrieves, and manages data in document format, typically JSON, BSON, or XML.
- Schema on Read

Problem

- Lack of ACID Transactions.
- Limited Querying Capabilities:
- Eventual Consistency
- Data Duplication

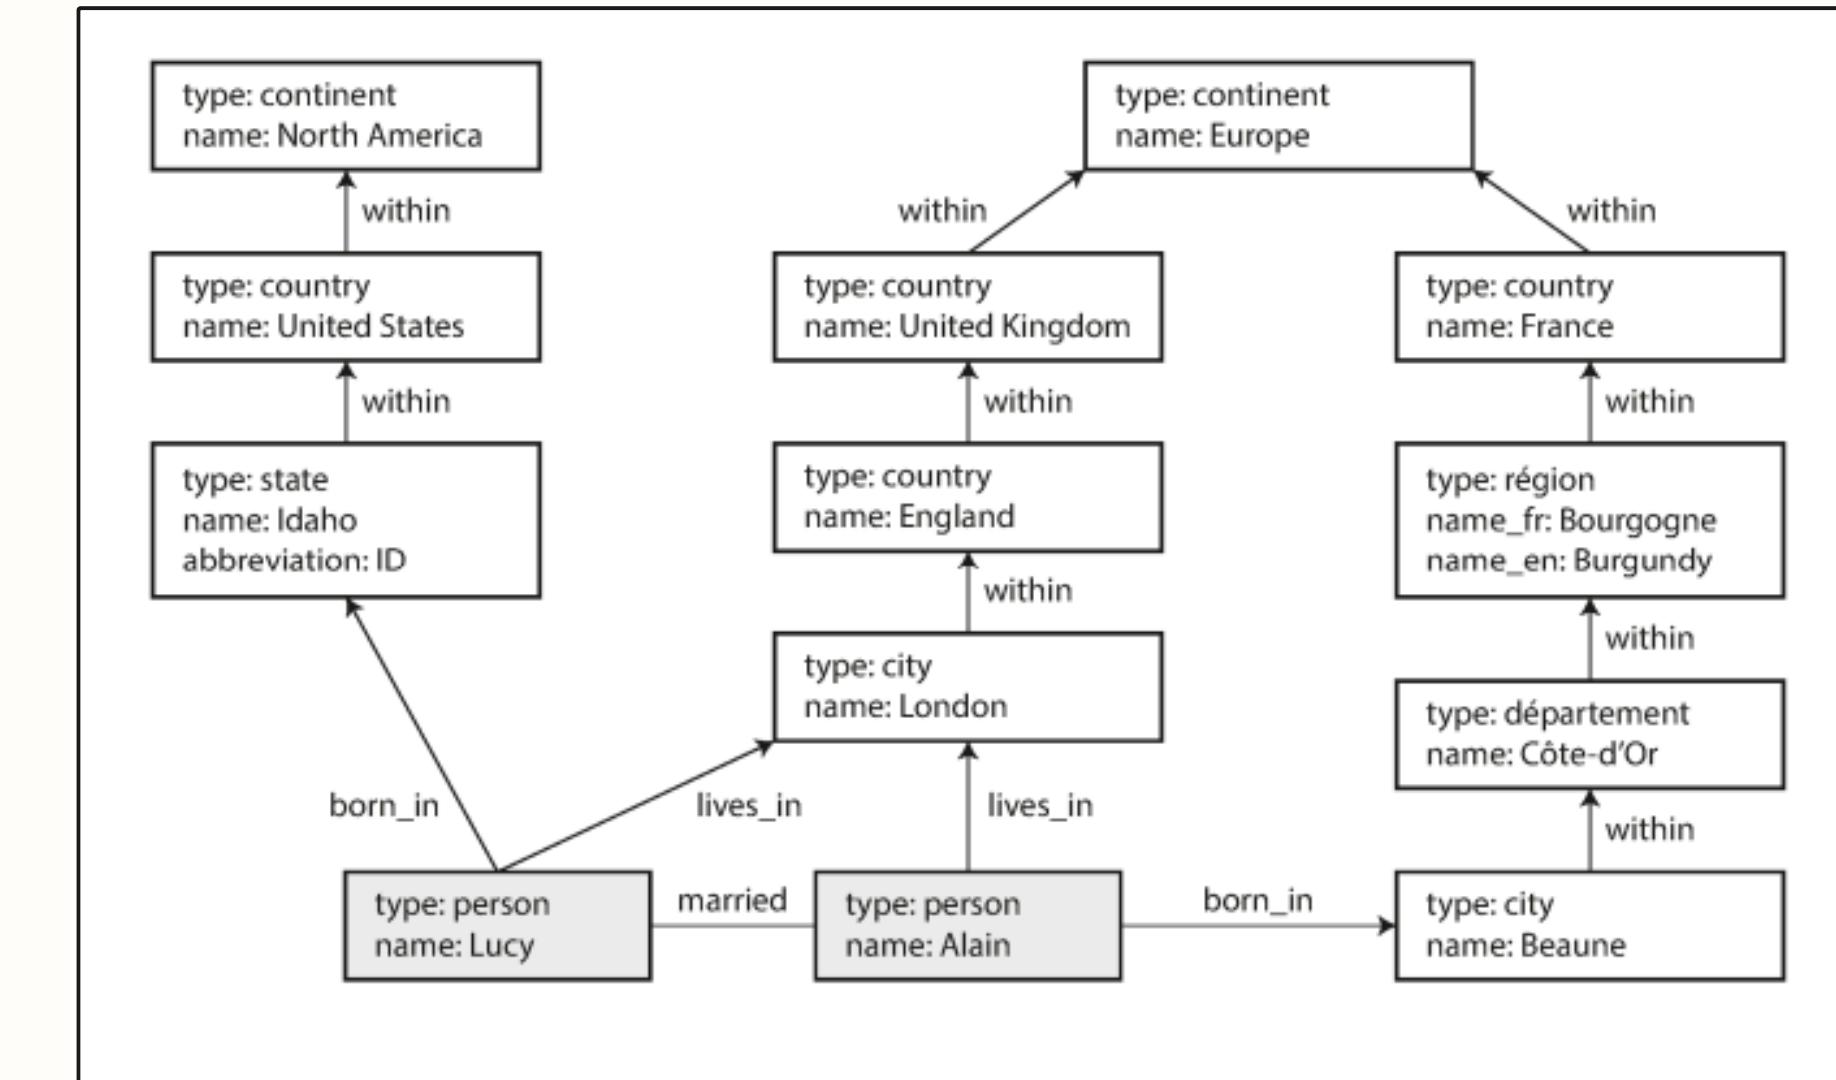
```
{  
  "user_id": 251,  
  "first_name": "Bill",  
  "last_name": "Gates",  
  "summary": "Co-chair of the Bill & Melinda Gates... Active blogger.",  
  "region_id": "us:91",  
  "industry_id": 131,  
  "photo_url": "/p/7/000/253/05b/308dd6e.jpg",  
  "positions": [  
    {"job_title": "Co-chair", "organization": "Bill & Melinda Gates Foundation"},  
    {"job_title": "Co-founder, Chairman", "organization": "Microsoft"}  
,  
  "education": [  
    {"school_name": "Harvard University", "start": 1973, "end": 1975},  
    {"school_name": "Lakeside School, Seattle", "start": null, "end": null}  
,  
  "contact_info": {  
    "blog": "http://thegatesnotes.com",  
    "twitter": "http://twitter.com/BillGates"  
  }  
}
```

Graph

- A graph database represents data as nodes (entities) and edges (relationships), making it ideal for handling complex, highly connected data.

Characteristic

- Relationship-Centric
- Fast Traversal
- Flexible Schema
- Rich Metadata



Comparison

Relational	Document	Graph
Better joins	Schema flexibility	Best for complex, interconnected relationships
Better many to one and many to many relationship	Better performance due to locality	Efficient traversal of complex relationships
Schema on write	Schema on read	Flexible schema
Requires rewriting the entire row	Updates a document without changing the entire content	Updates nodes independently
Reads only the necessary data	Requires reading the entire document for retrieval	Quick access to related data through graph traversals
Performance may degrade with complex queries	Performance may degrade with complex queries	Performance remains stable for complex relationship queries



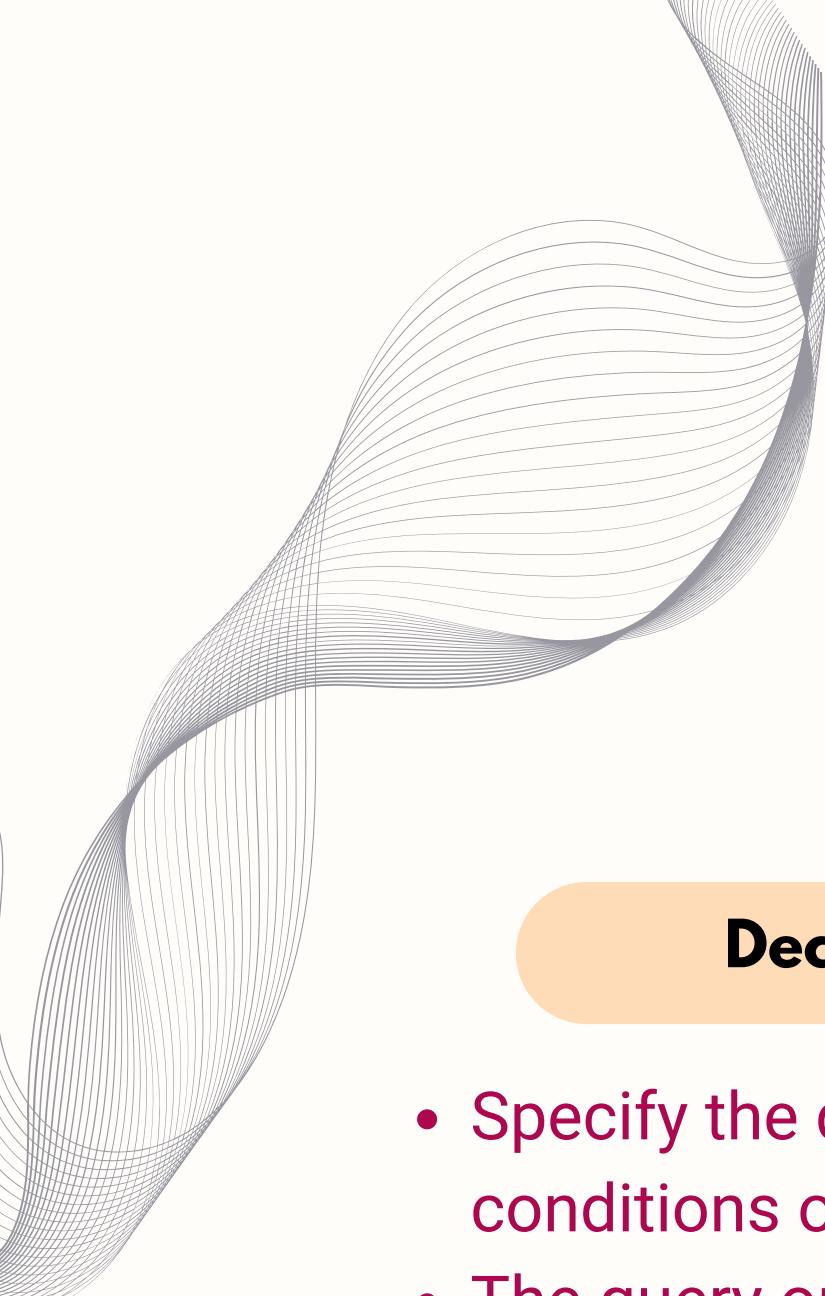
Query Languages

Declarative

```
Select * from animals where family='sharks';
```

Imperative

```
function getSharks() {  
    var sharks = [];  
    for(var i=0;i<animals.length;i++){  
        if(animals[i].family === "Sharks"){  
            sharks.push(animals[i]);  
        }  
    }  
    return sharks;  
}
```



Query Languages

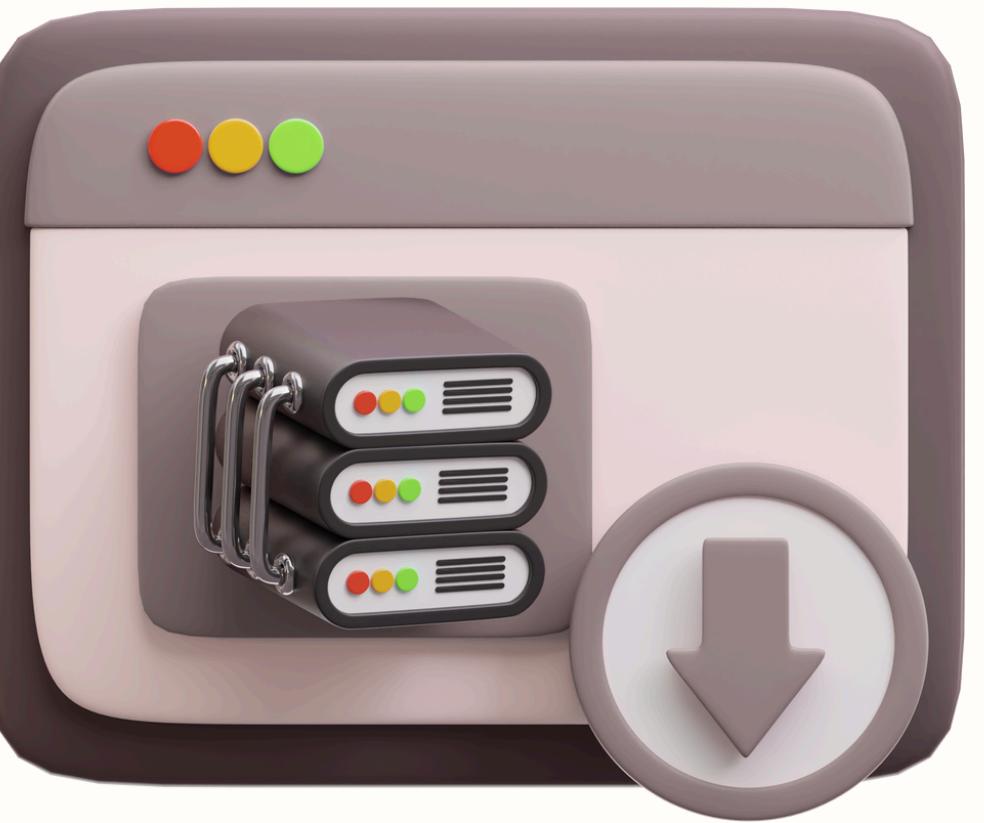
Declarative

- Specify the desired data pattern (like conditions or transformations).
- The query optimizer finds the best way to retrieve matching data.
- Easier to parallelize, as execution can be handled automatically.

Imperative

- Much more verbose but flexible. Coding the specific operations
- Parallel code harder to write for each query
- Tells the computer to perform certain operation in a certain order.
- Hard to parallelize because it specifies instructions to be executed in a particular order

Chapter 3



Storage and Retrieval

What does a database do?



Store data



**Give data when
requested**

As a Developer it is essential for selecting the right storage engine for your application's workload, optimizing performance, and preventing bottlenecks as usage scales.

Data Structures

Log Structure

- SS Tables
- LSM Trees

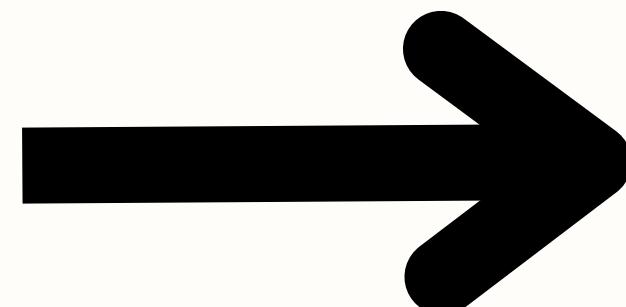
Index Based

- B-Trees
- Hash

Why Index?

```
$ db_set 42 '{"name":"San Francisco","attractions":["Exploratorium"]}'  
  
$ db_get 42  
{"name":"San Francisco","attractions":["Exploratorium"]}  
  
$ cat database  
123456,{"name":"London","attractions":["Big Ben","London Eye"]}  
42,{"name":"San Francisco","attractions":["Golden Gate Bridge"]}  
42,{"name":"San Francisco","attractions":["Exploratorium"]}
```

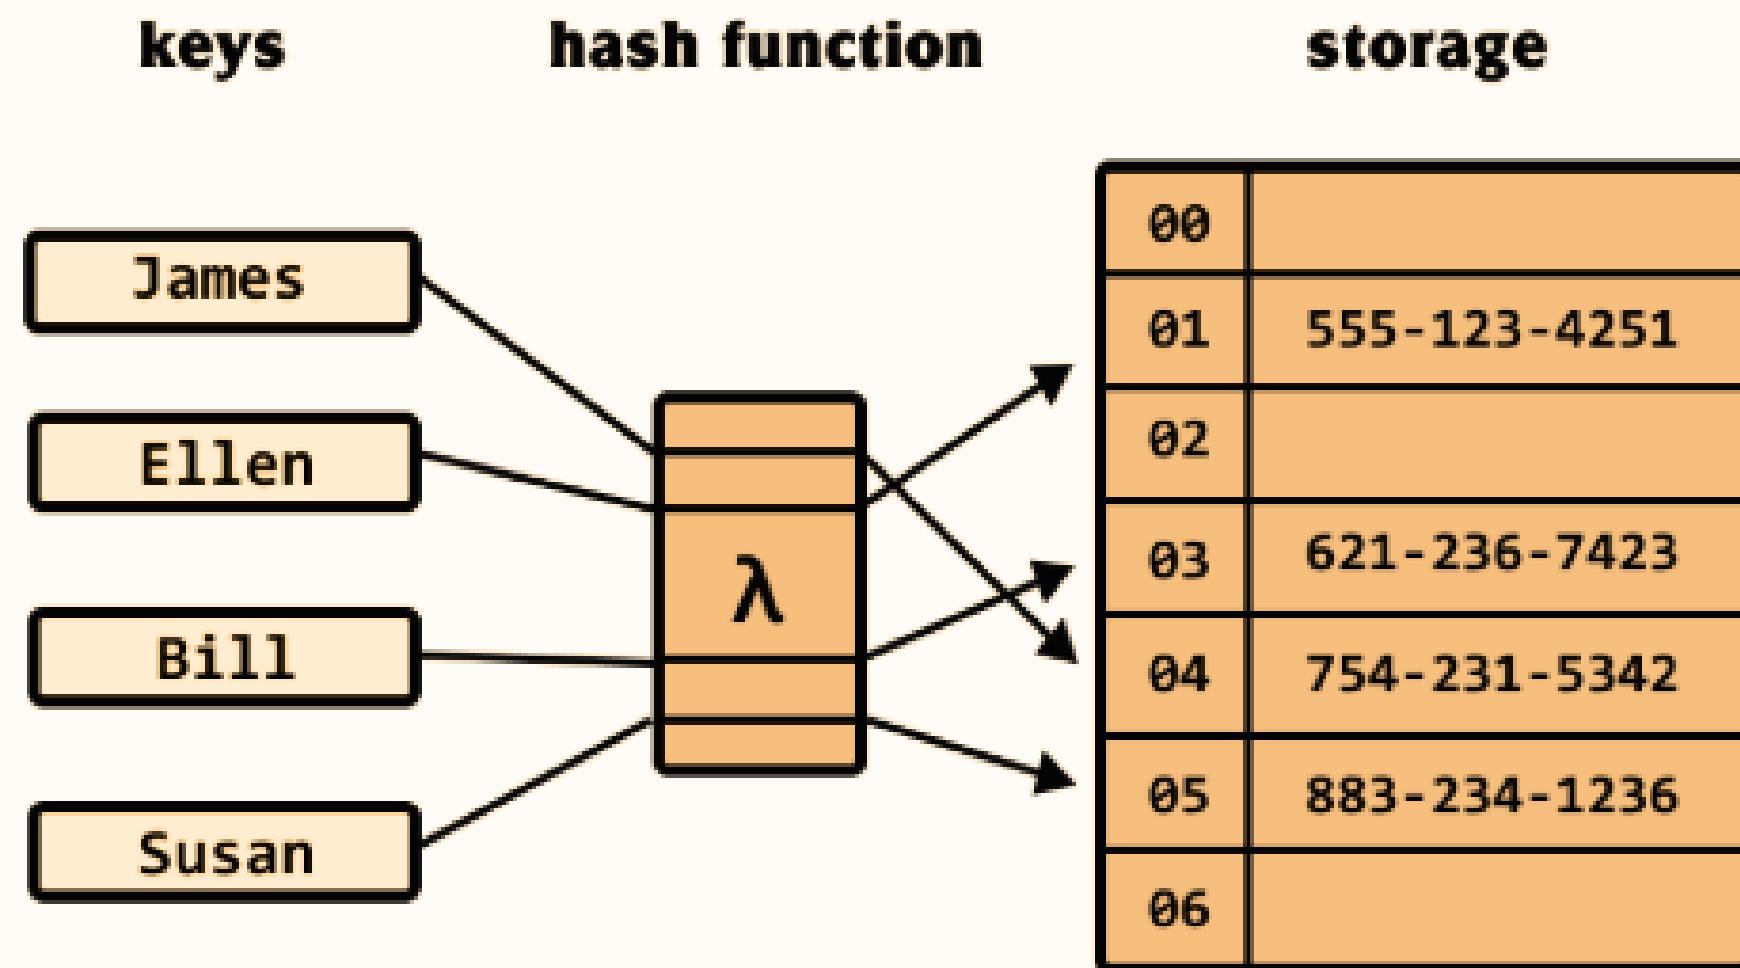
- Inefficient Lookups ($O(n)$)
- Redundant Data Growth



- Optimized Search and Updates with Indexing

Index Based

Hash



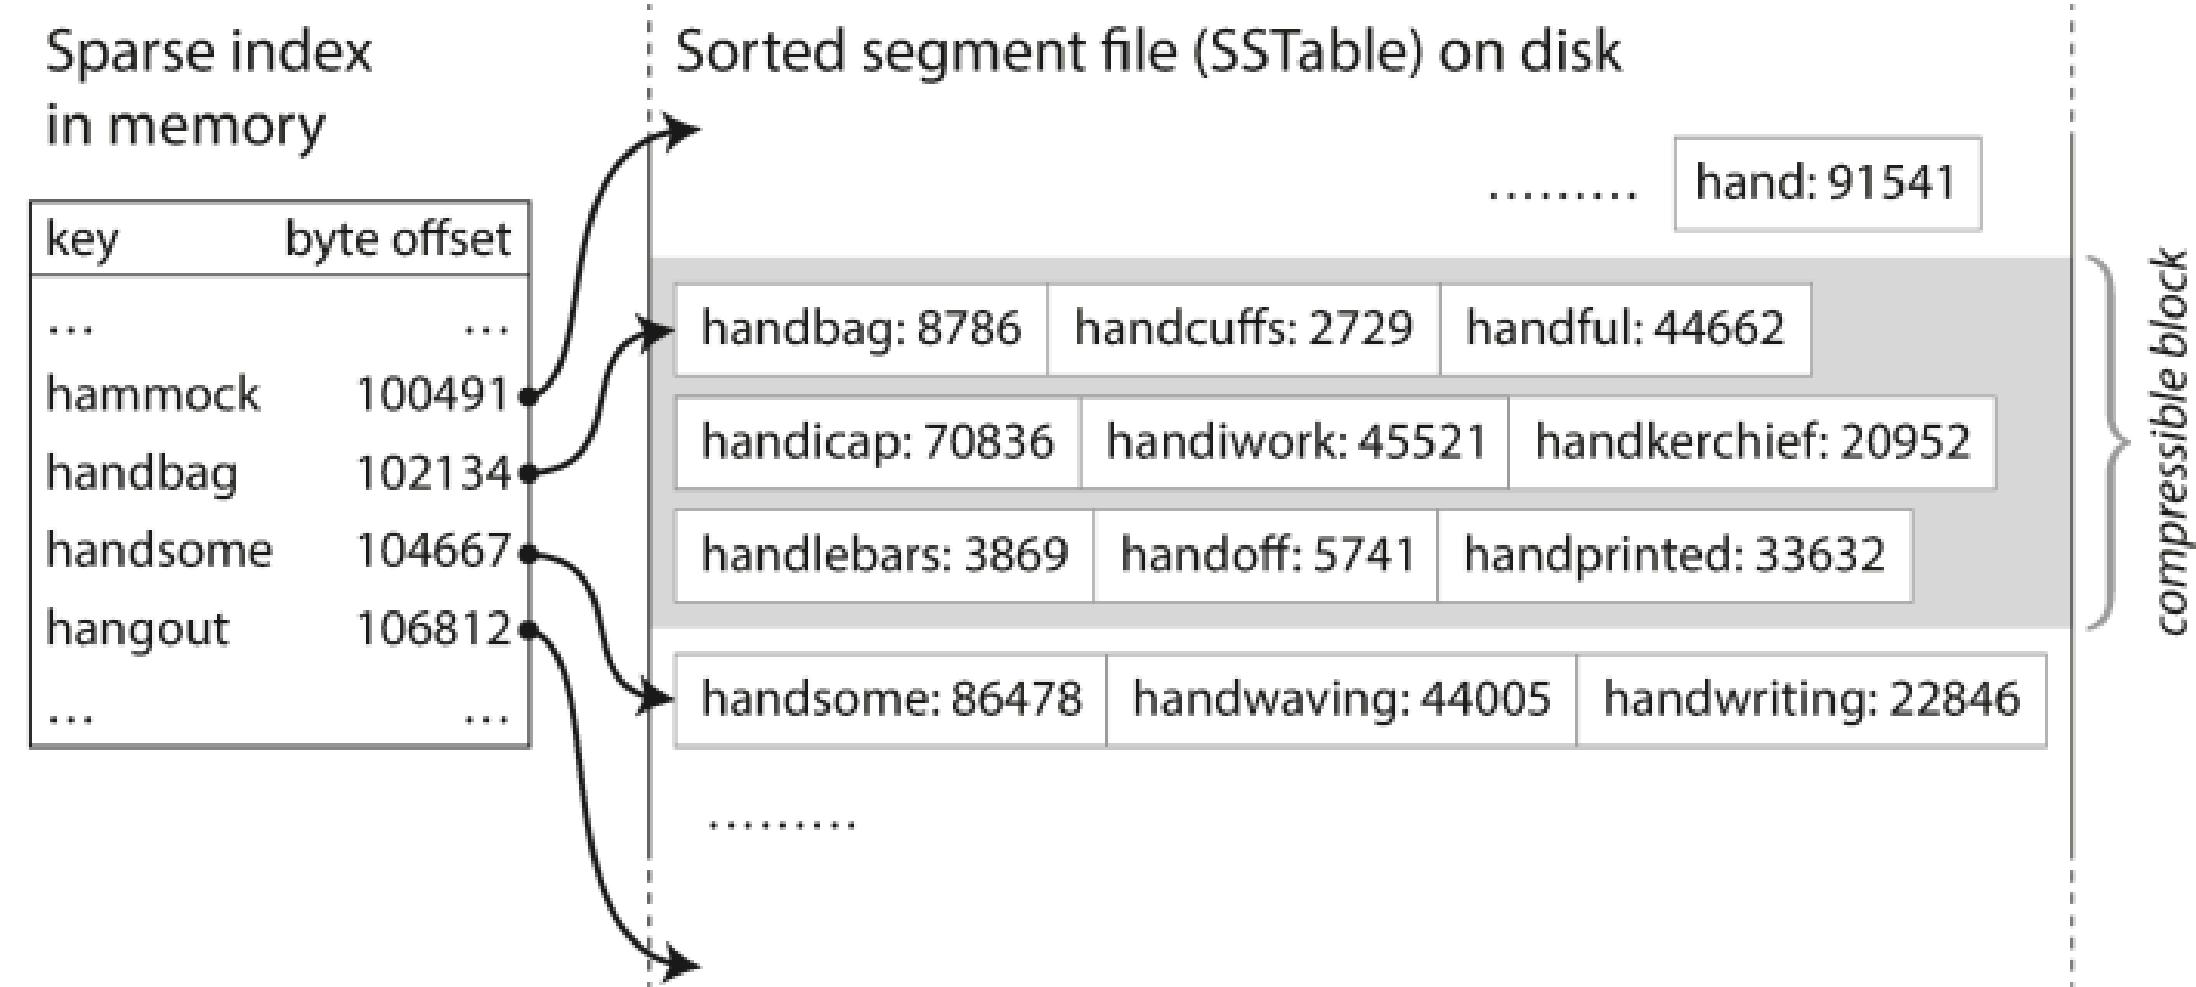
Hash Map

Limitation

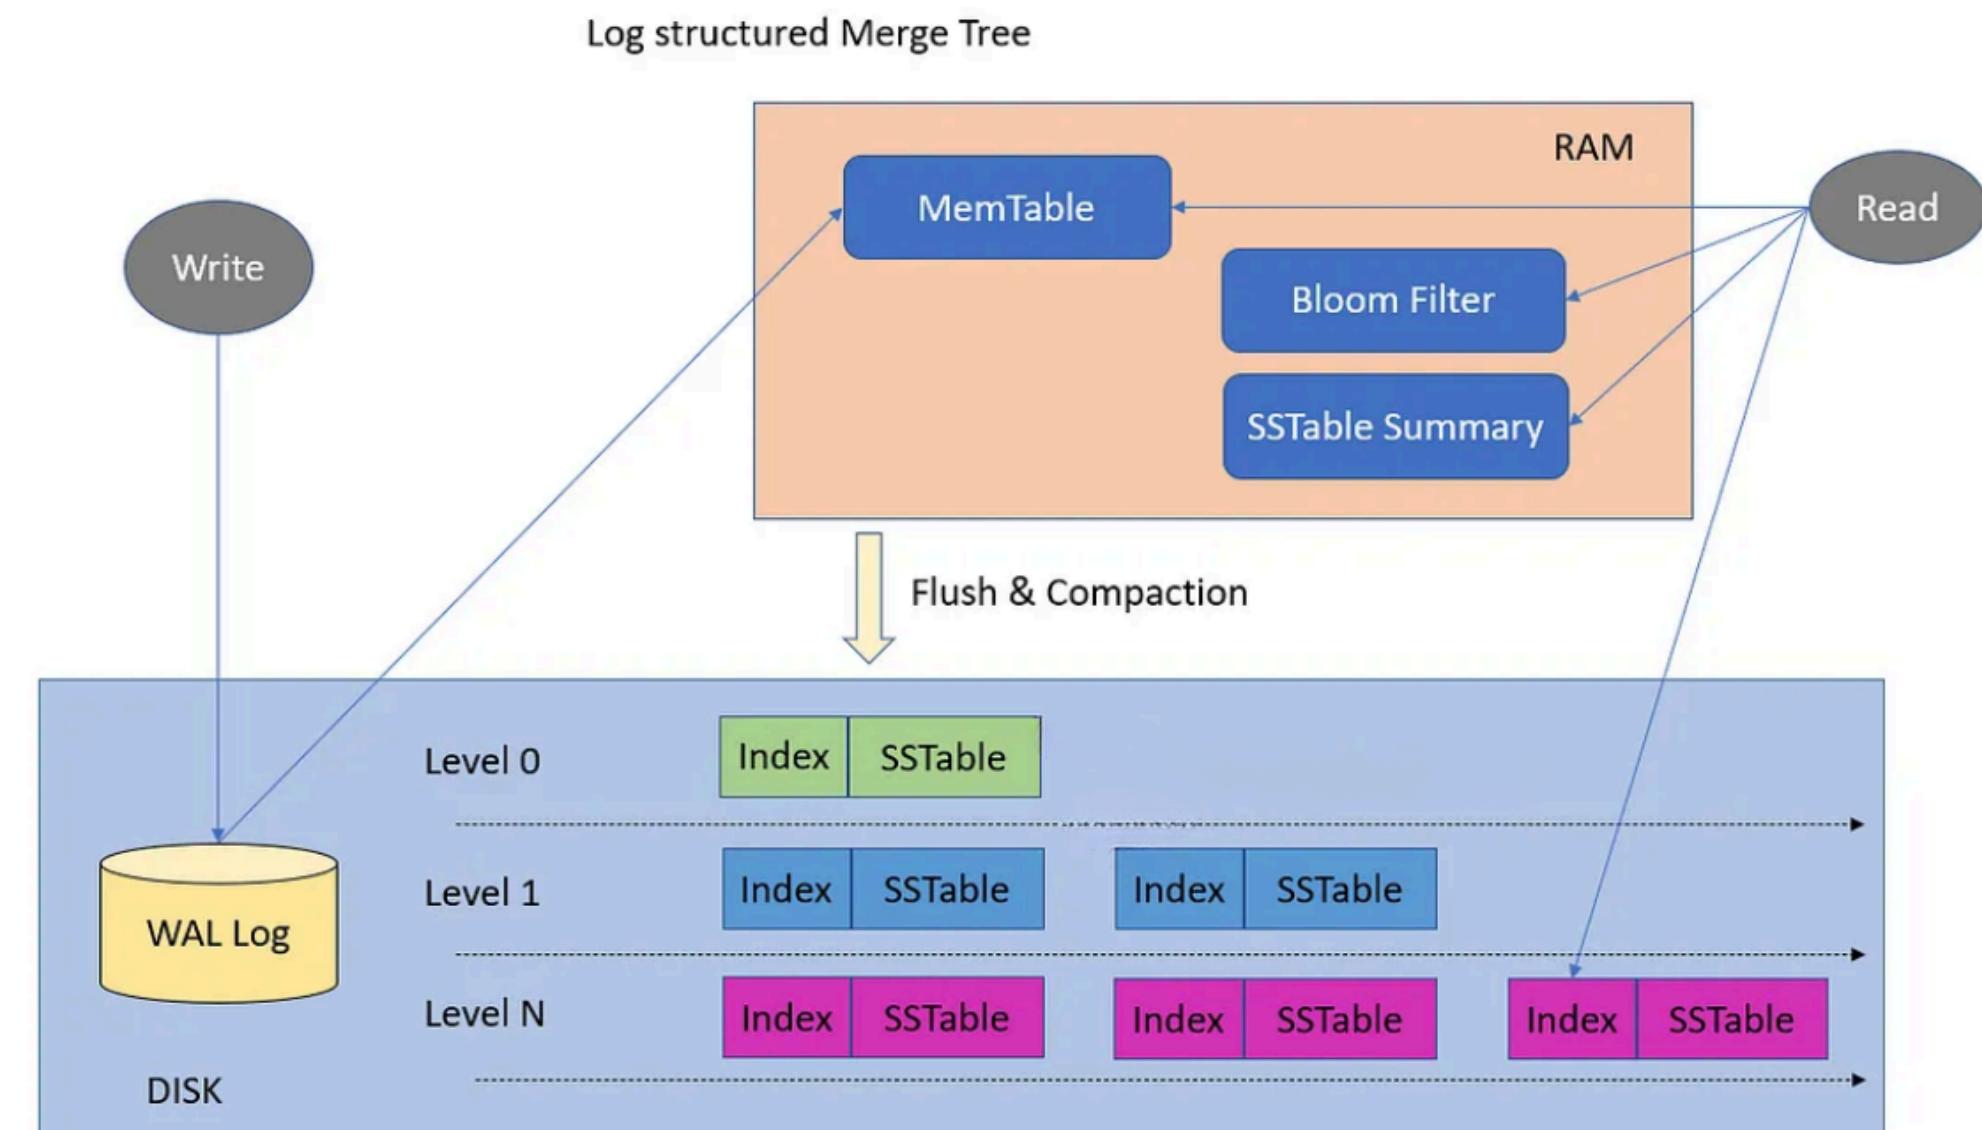
- Memory and Disk Access Limitation – The hash table must fit entirely in RAM, and if it exceeds available memory, storing it on disk leads to excessive random access, significantly slowing down operations.
- Inefficient for Range Queries – Hash tables do not support range-based lookups, requiring individual key lookups for queries within a specific range
- Collisions - that can be handled with chaining or open addressing

Data Structures

SS Tables



LSM Trees



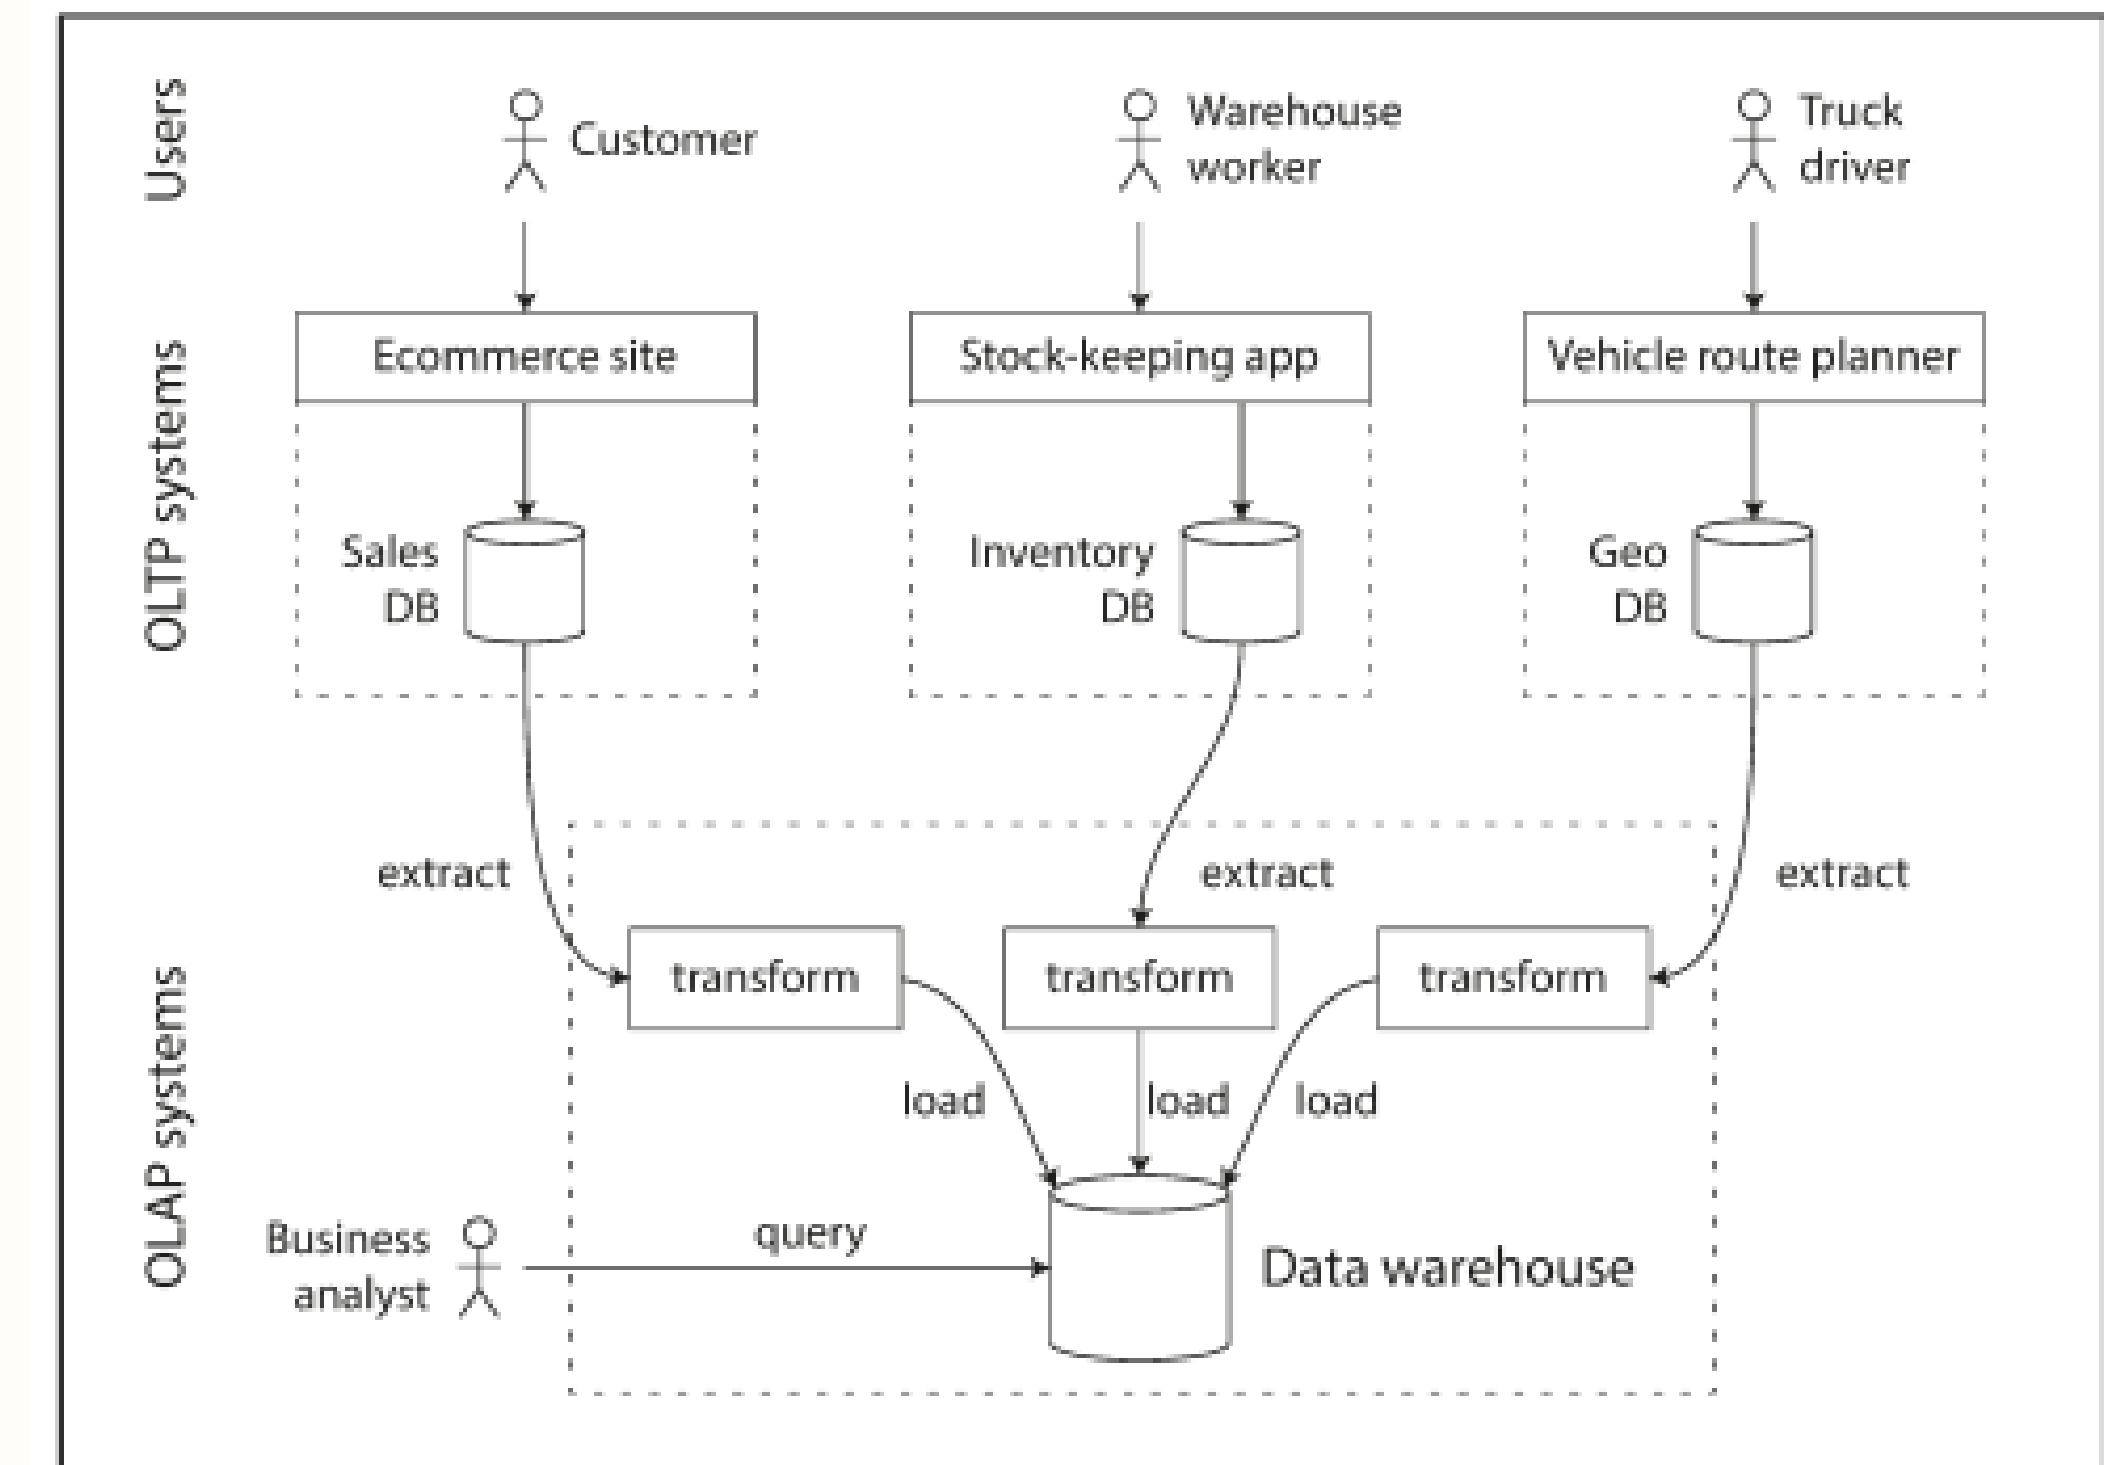


**Understanding data structures helps us design efficient databases,
but how data is accessed depends on whether we focus on
transactions or analytics.**

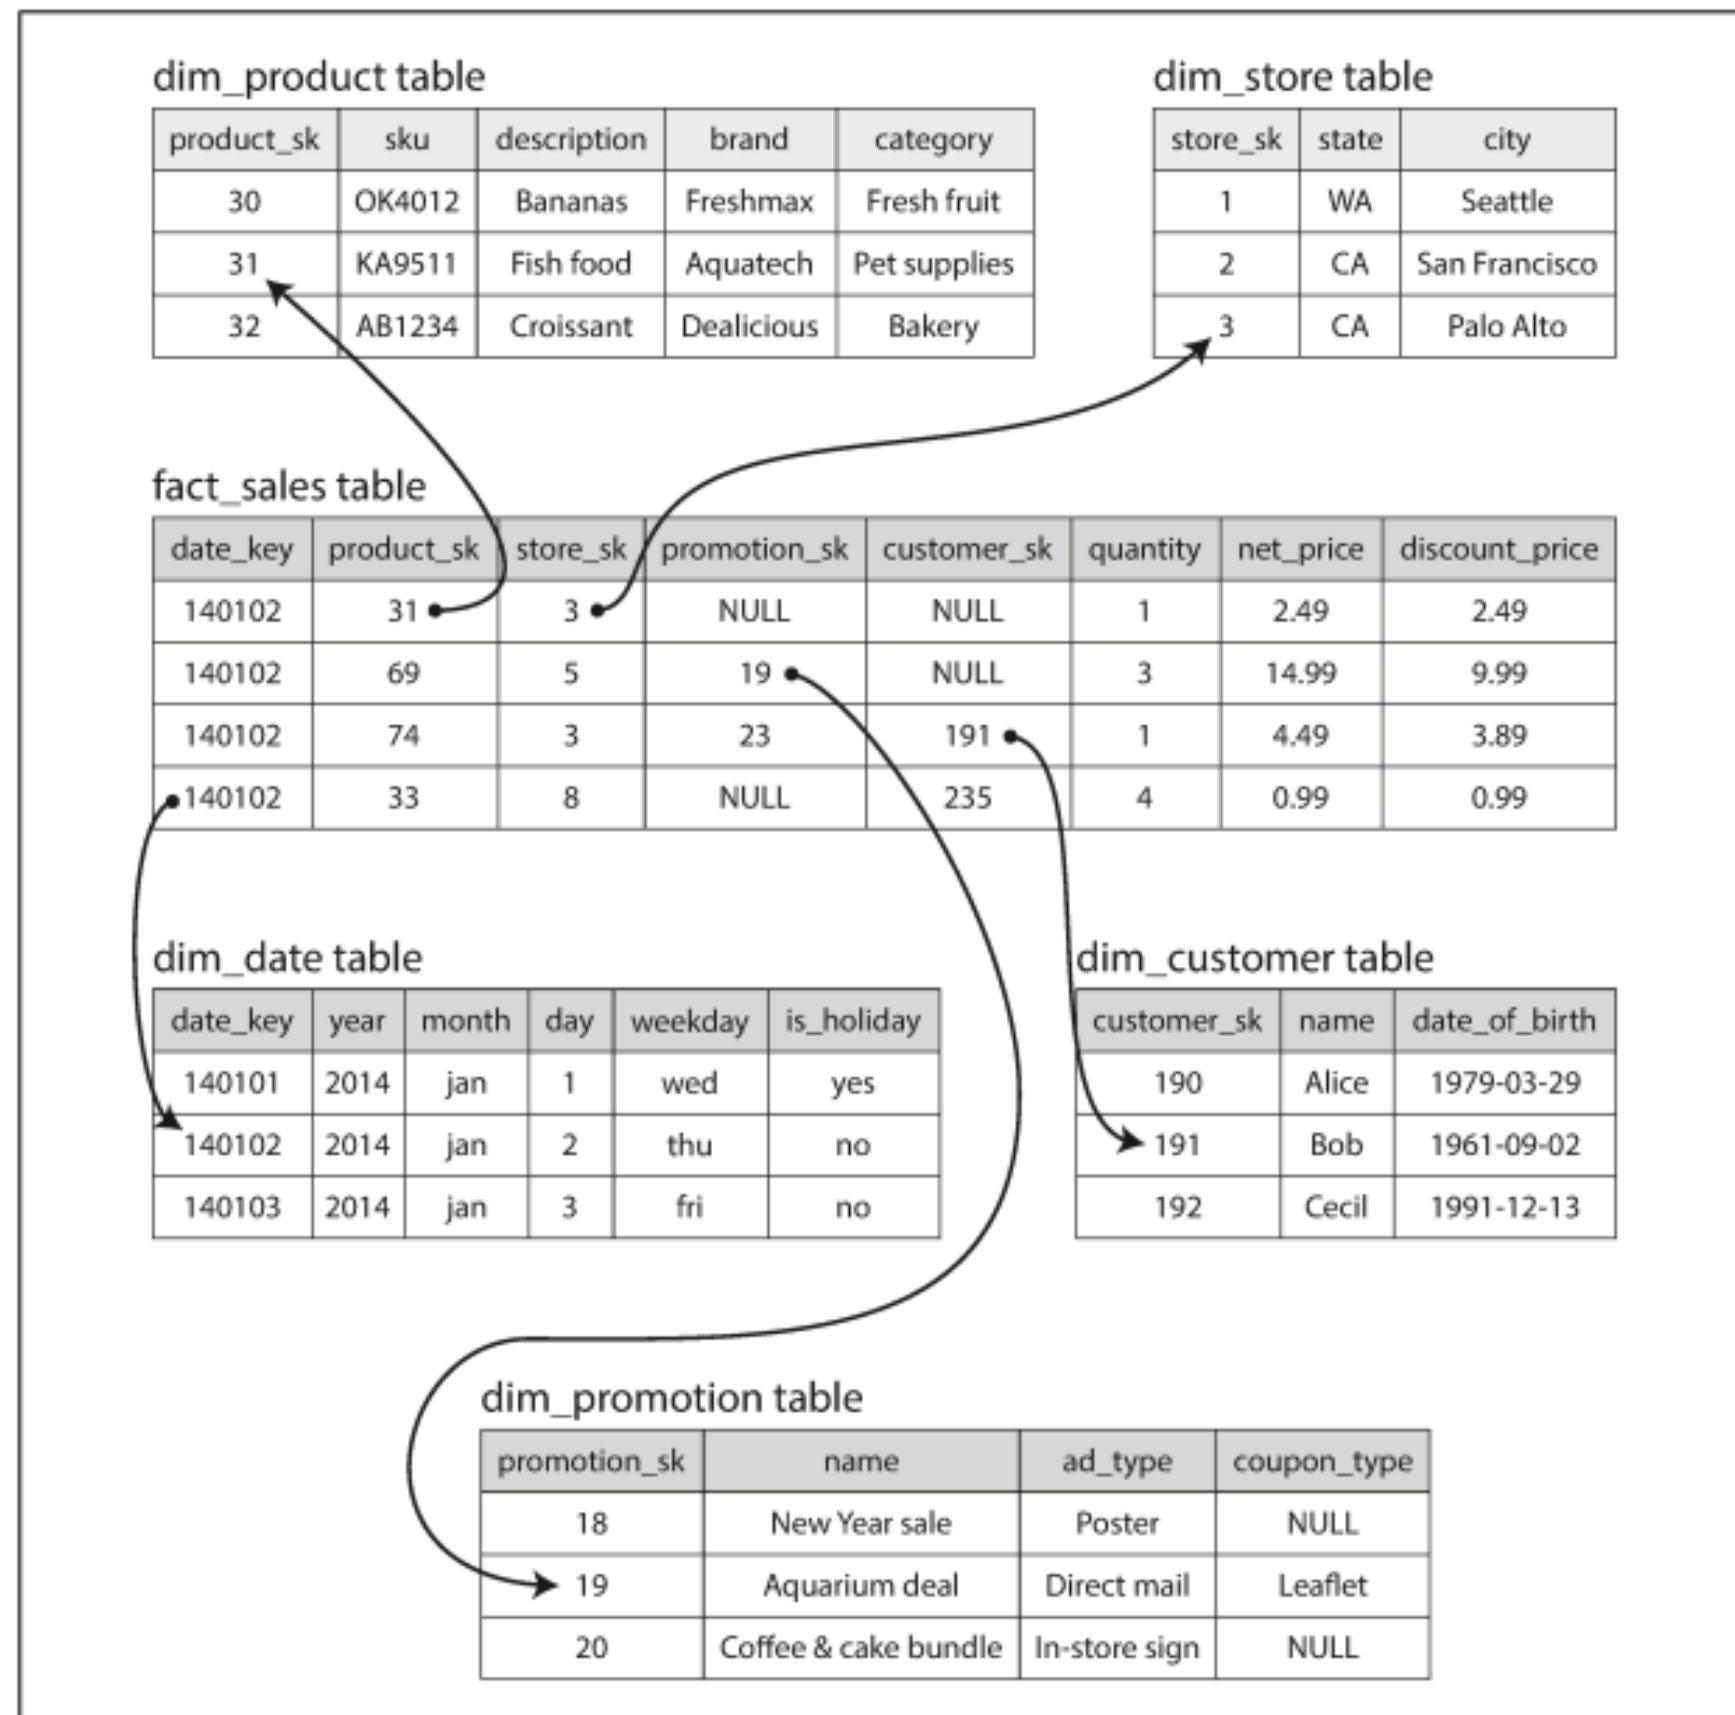
OLTP vs OLAP

Property	OLTP	OLAP
Main read pattern	Small number of reads per query, fetched by key.	Aggregate over large number of records.
Main write pattern	Random-access, low-latency writes from user input.	Bulk import (ETL) or event stream
Primarily used by	End user/customer, via web application	Internal analyst, for decision support.
What data represents	Latest state of data.	History of events that happened over time.
Dataset Size	Gigabytes to terabytes.	Terabytes to petabytes.

Extract, Transform Load(ETL) to Data Warehouse



Star Scheme in Data Warehouse



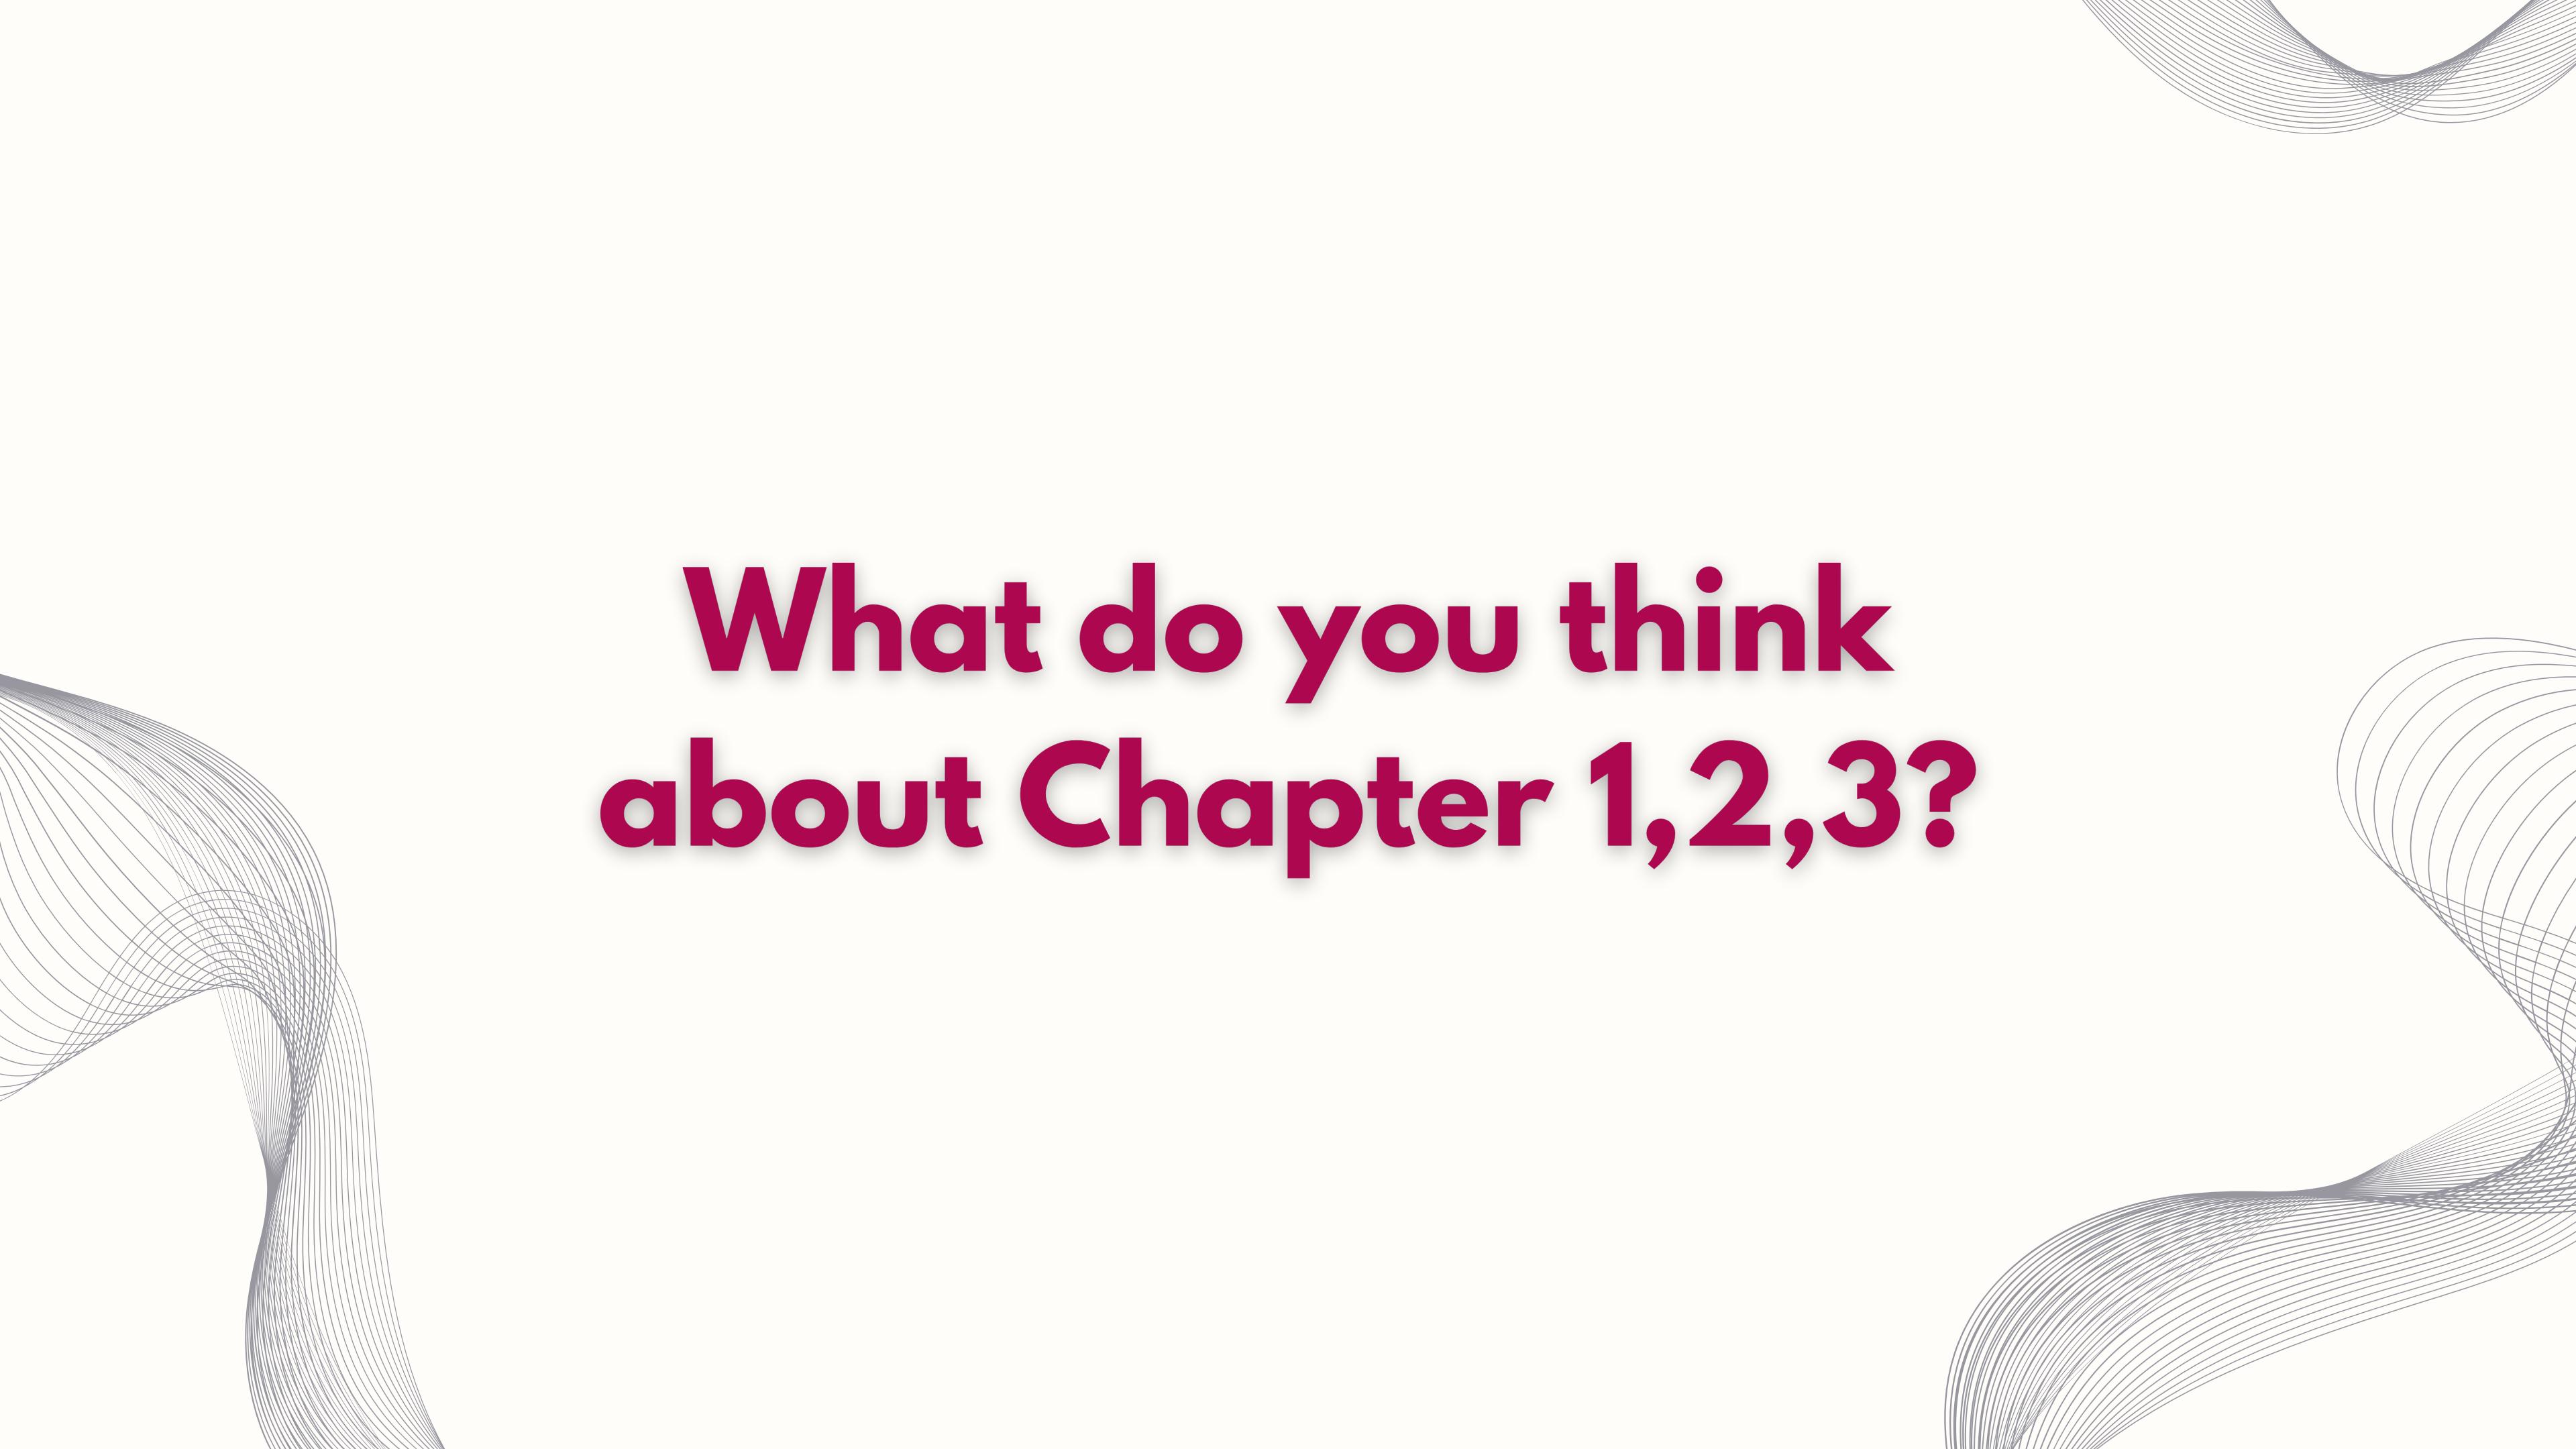
Column Oriented Storage

What?

It is a method of storing data where all the values in a single column are stored together, instead of storing entire rows together.

Why?

- Row storage is fast for CRUD but inefficient for OLAP.
- Column storage reads only needed columns, reducing I/O and speeding up queries.



**What do you think
about Chapter 1,2,3?**

Thank You So Much

